

Entropy Maximization in Sparse Matrix by Vector Multiplication ($\max_E \text{SpMV}$)

ABHISHEK JAIN, ISMAIL BUSTANY, HENRI FRAISSE, MANSIMRAN BENIPAL, DINESH GAITONDE, and PAOLO D'ALBERTO, Xilinx

The peak performance of any SpMV depends primarily on the available memory bandwidth and the capability to use it effectively. GPUs, ASICs, and new FPGAs have higher and higher bandwidth; however, for large scale and highly sparse matrices we find still difficult utilizing this bandwidth because the SpMV random access pattern and workload imbalance. We propose a matrix permutation pre-processing step that aims to maximize the entropy of the distribution of the nonzero elements. We seek any permutation that uniformly distributes the non-zero elements' distribution, thereby generating a SpMV problem that is amenable to work load balancing or to speed up sort algorithms. We conjecture these permutations would be most effective for matrices with no dense rows or columns and, as in preconditioning, when the matrix is reused. We shall show that entropy maximization is an optimization that any architecture may take advantage although in different ways. Most importantly, any developer can consider and deploy. We shall present cases where we can improve performance by 15% on AMD-based systems.

ACM Reference Format:

Abhishek Jain, Ismail Bustany, Henri Fraisse, Mansimran Benipal, Dinesh Gaitonde, and Paolo D'Alberto. 2020. Entropy Maximization in Sparse Matrix by Vector Multiplication ($\max_E \text{SpMV}$). 1, 1 (August 2020), 27 pages.

1 INTRODUCTION

To define the scope of this work, the obvious questions to ask are: first, what randomization or entropy maximization is in the context of sparse matrices; second, why would we use it; third, when it does work. We shall provide formal definitions in the following sections. Briefly, we will permute randomly the rows and columns of a sparse matrix before multiplying it with a dense vector (SpMV) with the aim of speeding this operation. Undoubtedly, this scheme requires some restrictions about the matrix structure, one among them is that it has no or few dense columns or rows. In the case, where there are dense columns or rows, a sparse/dense partitioning scheme should be used. For the remainder of this manuscript, we shall assume the former nonzero structure. We use randomization because it is the poor man's way for preconditioning SpMV in our context, and we do not mean it in a pejorative sense.

Preconditioning speeds up the convergence rate of an iterative linear solver by linearly transforming the associated matrix into a form that affords a faster reduction of the residual error at every iteration. The cost of this transformation is justified by the runtime reduction it affords. Likewise, we foresee randomization playing a similar role for SpMV in the context of iterative linear solvers and other methods (e.g. in convolutions) where the matrix is reused.

Sparse linear algebra and GraphBLAS kernels are memory bound and there is a common thread in the scientific computing community to develop acceleration libraries mostly for multi-core systems. These predominantly include multi-core processors and GPUs. The goal is a balanced work distribution and, when applicable, minimal communication

Authors' address: Abhishek Jain, abhishek@xilinx.com; Ismail Bustany, ismailb@xilinx.com; Henri Fraisse, henrif@xilinx.com; Mansimran Benipal, mansimra@xilinx.com; Dinesh Gaitonde, dineshg@xilinx.com; Paolo D'Alberto, paolod@xilinx.com, Xilinx, 2100 Logic Dr, San Jose, California, 95124.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2020 Association for Computing Machinery.

Manuscript submitted to ACM

[6, 11]. When storage strategy and algorithms must be considered together then GPUs provide the work horse for abundant thrust in research [1]. These works aim at optimal solutions and strive for a clear and complete understanding/exploitation of the software-hardware interface; usually the hardware is composed of symmetric computational units. Interestingly, the SpMV's space and time complexity, which are small, may not warrant more performance because we typically end up utilizing only one-thousandth fraction of the available hardware capacity.

The peak performance of any SpMV accelerator depends primarily on the available memory bandwidth (i.e., DRAM such as DDR or HBM) and the capability of the accelerator to effectively use it. Because SpMV is memory-bound, a more important metric than peak performance alone is the fraction of bandwidth utilized, which captures the overall efficiency of the architecture. GPU platforms exhibit very high bandwidth, see the experimental Section 8: Ellesmere DDR5 224GB/s, Fiji HBM 512GB/s, and Vega 20 HBM 1TB/s. Although utilizing this much bandwidth efficiently is difficult for large scale and highly sparse matrices due to very high random access pattern. Custom architectures based on FPGA or ASIC devices can maximize bandwidth utilization by highly customized data-paths and memory hierarchy designs [3, 4, 13]. Most of the existing accelerators saturate the relatively low memory bandwidth available on FPGA platforms (less than 80 GB/s) [3, 4, 8, 10, 12, 13]. Modern FPGA platforms have multiple HBM stacks to provide large memory bandwidth. However, there is no implementation (currently available) that saturates all of the available DRAM bandwidth for SpMV kernel on HBM-enabled FPGA platforms. Scalability of accelerator design remains a major concern, and it is an active area of research.

FPGA platforms used in early works exhibit low peak performance due to the scarcity of external memory bandwidth [3, 7, 14]. For example, Microsoft's implementation of SpMV uses an FPGA platform which only has 2 DDR2-400 memory banks with a resulting bandwidth of 6.4 GB/s [7]. The accelerator is running at 100 MHz, it reads 64 Bytes of data every cycle, which corresponds to 5 non-zeros at every cycle (a non-zero is about 12 Bytes). At best, the peak performance is 10 double precision operations every cycle at 100 MHz, which is 1 GFLOPS (only). In 2009, Convey systems Inc. released the Convey HC-1 FPGA platform. It has 16 DDR2-677 memories resulting in overall 80 GB/s memory bandwidth [10]. The accelerator logic runs at 150 MHz. It consumes 512 Bytes of data every cycle, which corresponds to around 40 non-zeros every cycle. At best, the peak performance is 80 double precision operations every cycle at 150 MHz, which is 12 GFLOPS.

One of the key building blocks for custom architecture solutions is a multi-ported buffer used to storing vector entries [3]. During execution, multiple column indices are used as addresses to read corresponding vector entries; we shall provide more details about the application in Section 2. Designing a buffer with a very large number of read ports is challenging. One solution is *banking* as a mechanism to store partitioned vector entries. Although banking could allow very high throughput indexing unless the same entry is required multiple times and its reads are purely sequential causing loss of bandwidth. For example, hashing techniques and data duplication are possible solutions for this problem. However, another issue arises: When we distribute SpMV computations across p -nodes, some of the nodes, say k , finish later than the rest because of unbalanced work loads (i.e., number of nonzero element) in row/column major traversal. This is a common phenomena for matrices where few rows or columns are dense. These k nodes are referred to as *laggard nodes*. By applying random permutation of columns/rows, we are attempting to balance the loads across all p workers so that there are no laggards. From this hardware vantage point, randomization or maximizing the entropy of the non-zero element distribution is an optimization transform and provides a clear context for our work.

Clearly, optimally accelerating SpMV is a hard many-parameters optimization problem dependent on the choice of algorithm, data structures, and dedicated hardware (CPU, GPUs, FPGA's, Custom ASIC's). Rather, our goal is to provide a tool, we may say a naive tool, to help understand how the structure of the matrix may affect the HW-SW

solution. For the readers in the field of algorithms, SpMV can be mapped into a sorting algorithm. For example, finding elements $x_{i,j}$ and $x_{i,k>j}$ in a sparse matrix requires to find row i and then columns j and k . Sorting is a method to find if an element is in a list with no prior or limited knowledge of its contents. Sorting can be used to prepare the matrix and to find elements in between sparse matrices and sparse vectors. In custom architectures, sorting networks are used to route matrix and vector elements to functional units. In a sense, if one is stuck with a sorting algorithm and a poor distribution, randomization may alter the distribution and throttle performance. Interestingly, the best sorting algorithm is a function of the distribution of the elements [5, 9].

We organize this work as follows: In Section 2, we define the matrix by vector operation; in Section 3, we define what we mean by randomization or entropy maximization. We use randomization to create a uniform distribution in Section 5 and measure uniformity by entropy in Section 4. We present how we drive our experiments to show the effects of randomization in Section 6. In the last sections, we present a summary of the results: we present our work loads for the given benchmarks in Section 7, and the complete set of measures for an AMD CPU and GPUs systems in Section 8.

2 BASIC NOTATIONS

Let us start by describing the basic notations so we can clear the obvious (or not). A Sparse-matrix by vector multiplication SpMV on an (semi) ring based on the operations $(+, *)$ is defined as $\mathbf{y} = \mathbb{M}\mathbf{x}$ so that $y_i = \sum_j M_{i,j} * x_j$ where $M_{i,j}=0$ are not represented nor stored. Most of the experimental results in Section 8 are based on the classic addition $(+)$ and multiplication $(*)$ in floating point precision using 64 bits (i.e., double floating point precision) albeit are extensible to other semi-rings. For instance, it is well known that SpMV defined on the semi-ring $(\min, +)$ is a kernel in computing an all-pairs shortest paths starting with a graph adjacency matrix, and in using a Boolean algebra we can check if two nodes are connected, which is slightly simpler.

We identify a sparse matrix \mathbb{M} of size $M \times N$ as having $O(M + N)$ non-zero elements, number of non zero nnz . Thus the complexity of $\mathbb{M}\mathbf{x}$ is $O(M + N) \approx 2nnz$. Also, we must read at least nnz elements and thus the complexity is $\Theta(M + N) \approx nnz$. We can appreciate that reading the data is as complex as the overall operation. Of course, the definition of sparsity may vary. We represent the matrix \mathbb{M} by using the coordinate list COO or and the compressed sparse row CSR¹ formats. The COO represents the non-zero of a matrix by a triplet (i, j, v) ; very often there are three identical-in-size vectors for the ROW, COLUMN, and VALUE. The COO format takes $3 \times nnz$ space and two consecutive elements in the value array are not bound to be neither in the same row nor column. In fact, we know only that $VALUE[i] = M_{ROW[i], COLUMN[i]}$.

The CSR format stores elements in the same row and with increasing column values consecutively. There are three arrays V, COL, and ROW. The ROW is sorted in increasing order. Its size is M , and $ROW[i]$ is an index in V and COL describing where i -th row starts (i.e., if row i exists). Accordingly, $M_{i,*}$ is stored in $V[ROW[i] : ROW[i + 1]]$. The column indices are stored at $COL[ROW[i] : ROW[i + 1]]$ and sorted increasingly. The CSR format takes $2 \times nnz + M$ space and a row vector of the matrix can be found in $O(1)$.

The computation $y_i = \sum_j M_{i,j} * x_j$ is a sequence of scalar products and, using the CSR format, is computed as follows:

$$Index = ROW[i] : ROW[i + 1]$$

$$y_i = \sum_{\ell \in Index} V[\ell] * x_{COL[\ell]}$$

¹a.k.a. Compressed row storage CRS.

The matrix row is contiguous (in memory) and rows are stored in increasing order. However, the access of the dense vector \mathbf{x} has no particular pattern, well increasing.

The COO format can be endowed with certain properties. For example, we can sort the array by row and add row information to achieve the same properties of CSR. In contrast, transposing a "sorted" COO matrix simply entails swapping of the arrays ROW and COL. Think about matrix multiply (one of us does constantly). Each scalar product achieves peak performance if the reads of the vector \mathbf{x} are streamlined as much as possible and so the reads of the vector V . If we have multiple cores, each could compute a subset of the y_i and a clean data load balancing can go a long way. If we have few functional units, we would like to have a constant stream of independent $*$ and $+$ operations but with data already in registers. That is, data pre-fetch will go a long way especially for $x_{COL[i]}$, which may have an irregular pattern.

3 RANDOMIZATION AND ENTROPY MAXIMIZATION

We define *Randomization* as row or column permutation transform of the matrix \mathbb{M} (thus a permutation of \mathbf{y} and \mathbf{x}), and we choose these by a pseudo-random process. The obvious question to ask is why should we seek randomization transform? The sparsity of a given matrix \mathbb{M} has a non-zero element distribution induced by the nature of the original problem or by some imposed ordering on the respective nodes of its associated graph. This distribution may be computationally incompatible with the chosen algorithm or architecture. For instance, it can induce some load imbalance in the computation. We could break this load imbalance by seeking to maximize entropy for this distribution. Our conjecture is that would favor the average case performance rather than the worse case when operating on the "max-entropy transformed" matrix.

For linear system solvers, if we know the matrix \mathbb{M} , and we know the architecture, preconditioning (when affordable) is a better solution. If we run experiments long enough, we choose the best permutation(s) for the architecture, permute \mathbb{M} , and go on testing the next. On one end, preconditioning exerts a full understanding of both the matrix (the problem) and how the final solution will be computed (architecture). On the other end, the simplicity of a random permutation requires no information about the matrix, the vector, and the architecture. Such a simplicity can be exploited directly in Hardware. We are after an understanding when randomization is just enough: We seek to let the hardware do its best with the least effort, or at least with the appearance to be effortless.

Interestingly, this work stems from a sincere surprise about randomization efficacy and its application on custom SpMV. Here, we wish to study this problem systematically so that to help future hardware designs. Intuitively, if we can achieve a uniform distribution of the rows of matrix \mathbb{M} we can have provable expectation of its load balancing across multiple cores. If we have a uniform distribution of accesses on \mathbf{x} we could exploit column load balancing and exploit better sorting algorithms: In practice, the reading of $x_{COL[i]}$ can be reduced to a sorting, and there we know that different sparsity may require different algorithms. This may be a lot to unpack but it translates to a better performance of the sequential algorithm without changing the algorithm or to improved bandwidth utilization.

We will show that (different) randomness affects architectures and algorithms differently, making randomization a suitable optimization transform especially when the application and hardware are at odds: Hardware (unless programmable) is difficult to change and the matrix sparsity is simple to change. We want to show that there is a randomness hierarchy that we can distinguish as global and local. There are simple-to-find cases where the sparsity breaks randomness optimization. For instance, matrices with dense rows or columns are better partitioned into sparse and dense components and operated on separately.

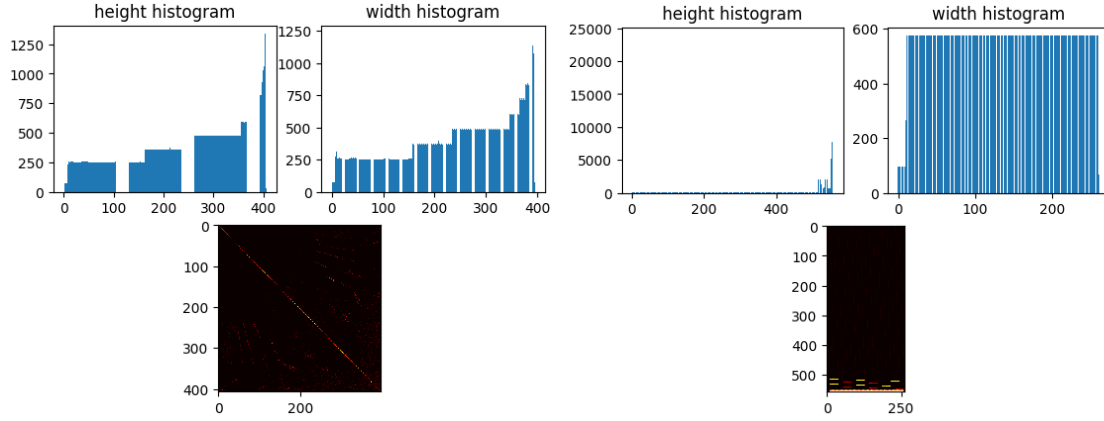


Fig. 1. Left: OPF 3754. Right: LP OSA 07. These are histograms where we represent normalized buckets and counts

4 ENTROPY

Patterns in sparse matrices are often visually pleasing, see Figure 1 where we present the height histogram, the width histograms, and a two-dimensional histogram as heat map. We will let someone else using AI picture classification. Intuitively, we would like to express a measure of uniform distribution and here we apply the basics: *Entropy*. Given an histogram $i \in [0, M - 1]$ $h_i \in \mathbb{N}$, we define $S = \sum_{i=0}^{M-1} h_i$ and thus we have a probability distribution function $p_i = \frac{h_i}{S}$. The *information* of bin i is defined as $I(i) = -\log_2 p_i$. If we say that the stochastic variable X has PDF p_i than the entropy of X is defined as.

$$H(x) = - \sum_{i=0}^{M-1} p_i \log_2 p_i = \sum_{i=0}^{M-1} p_i I(i) = E[I_x] \quad (1)$$

The maximum entropy is when $\forall i, p_i = p = \frac{1}{M}$; that is, we are observing a uniform distributed event. Our randomization should aim at higher entropy numbers. The entropy for matrix LP OSA 07 is 8.41 and for OPF 3754 is 8.39. We use the entropy specified in the Scipy stats module. A single number is concise and satisfying. If you are pondering why they are so close contrary to their sparsity we discuss this next.

5 UNIFORM DISTRIBUTION

We know that we should **not** compare the entropy numbers of two matrices because entropy does not use any information about the order of the buckets, it uses only their probabilities. By construction, the matrices are quite different in sparsity and in shapes, however their entropy numbers are close. Two matrices with the same number of non-zeros, spaced well enough in the proper number of bin, will have the same entropy. To appreciate their different sparsity, we should compare their entropy distributions by Jensen-Shannon measure [2] or we could use cumulative distribution function (CDF) measures, which imply an order. Here, we use a representation of a hierarchical 2D-entropy, see Figure 2, where the entropy is split into 2x2, 4x4 and 8x8 (or fewer if the distribution is not square). We have hierarchical entropy heat maps.

We can see that even a small 2D-entropy matrix summarizes the nature of the original matrix because it has spatial information. In this work, the entropy matrix is used mostly for visualization purpose more than for comparison purpose.

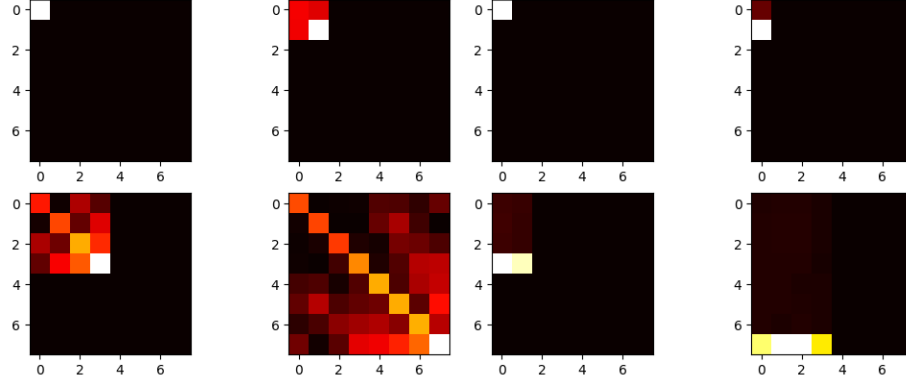


Fig. 2. Hierarchical 2D entropy for OPF 3754 (left) and LP OSA 07 (right).

Of course, we can appreciate how the matrix LP OSA 07 has a few very heavy rows and they are clustered. This matrix will help us showing how randomization need some tips. Now we apply row and column random permutation once by row and one by column: Figure 3: OPF has now entropy 11.27 and LP 9.26. The numerical difference is significant. The good news is that for entropy, being an expectation, we can use simple techniques like bootstrap to show that the difference is significant or we have shown that Jensen-Shannon can be used and a significance level is available. What we like to see is the the hierarchical entropy heat map is becoming *more* uniform for at least one of the matrix.

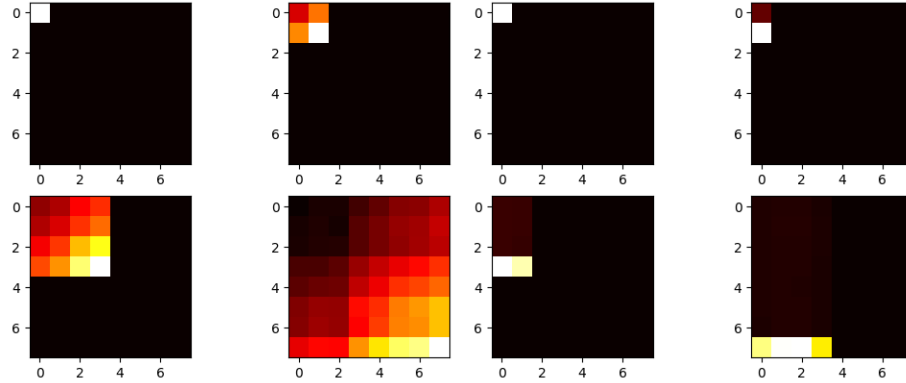


Fig. 3. Hierarchical 2D entropy after row and column random permutation for OPF 3754 (left) and LP OSA 07 (right).

In practice, permutations need some help especially for relatively large matrices. As you can see, the permutation affects locally the matrix. Of course, it depends on the implementation of the random permutation, we use *numpy* for this. It is reasonable that a slightly modified version of the original is still a random selection and unfortunately they seem too likely in practice. We need to compensate or help the randomization. If we are able to identify the row and column that divide high and low density, we could use them as pivot for a shuffle like in a quick-sort algorithm. We could apply a sorting algorithm but its complexity will the same of SpMV. We use a gradients operations to choose the element with maximum steepness, Figure 4 and 5.

LP achieves entropy 8.67 and 9.58 and OPF achieves 10.47 and 11.40.

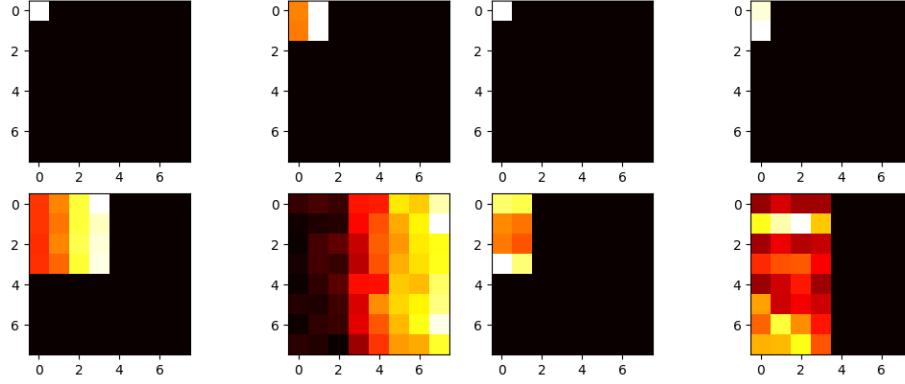


Fig. 4. Hierarchical 2D entropy after height gradient based shuffle and row random permutation for OPF 3754 (left) and LP OSA 07 (right).

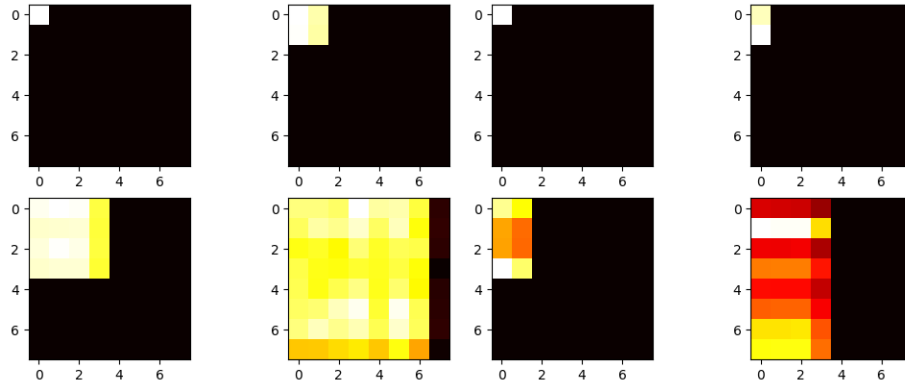


Fig. 5. Hierarchical 2D entropy after height and width gradient shuffle and row and column random permutation for OPF 3754 (left) and LP OSA 07 (right).

If the goal is to achieve a uniformly sparse matrix, it seems that we have the tools to compute and to measure such a sparsity. We admit that we do not try to find the best permutation. But our real goal is to create a work bench where randomization can be tested on different architectures and different algorithms. A randomization with a measurable uniform distribution is preferable than just random. We are interested to find out when random is enough or not enough. Also, consider that to achieve a uniform distribution, we do not need a random transformation and any permutation balancing the number of non-zero is possible, but for now not looked for.

6 MEASURING THE RANDOMIZATION EFFECTS

Whether or not this ever applied to the reader, when we have timed algorithms (i.e., measure execution time), we came to expect variation. The introduction of randomization may hide behind the ever present variance, after all these are

algorithms on *small* inputs: small error can be comparable to the overall execution time. Here, we must address this concern even before describing the experiments.

First, we execute every algorithm between 1000 and 5000 times. The time of each experiment is in the seconds, providing a granularity for which we are confident the measuring time error is under control. Thus, for each experiment we provide an average execution time: we measure the time and we divide by the number of trials. Cold starts, the first iteration, are still accounted. To make the measure portable across platform we present GFLOPS, that is, Giga (10^{12}) floating operations per second: $2 * nnz$ divided by the average time in seconds.

Then we repeat the same experiment 32 times. Permutations in *numpy* Python uses a seed that is time sensitive: thus every experiment is independent from the previous. The number 32 is an old statistic trick and it is a minimum number of independent trials to approximate a normal distribution. In practice, they are not but the number is sufficient for most of the cases and it is an excellent starting point.

A short hand legend: **Reg** is the regular matrix without any permutation; **R** stands for random Row permutation; **G-R** stands for gradient-based row shuffle and random row permutation; **G-C** stands for gradient-based column shuffle and random column permutation; **R-C** stands for random row and column permutation. This legend is used in the pictures to be concise, in the tables in the following sections, we use a verbose description. We shall clarify the gradient based approach in the experimental results section 8. Intuitively, we help the random permutation by a quick targeting of high and low volume of the histogram (and thus the matrix).

In Figure 6, we show two plots respectively of the CPU performance using COO and CSR SpMV algorithms for the matrix OPF 3754. The figure represents histograms: The x is GFLOPS and the y label is the number of counts. Thus we show what is the performance distribution of an algorithm. We can see that the CSR algorithms are consistent and the Regular (i.e., the original) has always the best performance. Also the variance of the computation time is small and the shape is approximately Gaussian. Different story for the COO, the permutations introduce long tails, thus 2× performance advantage.

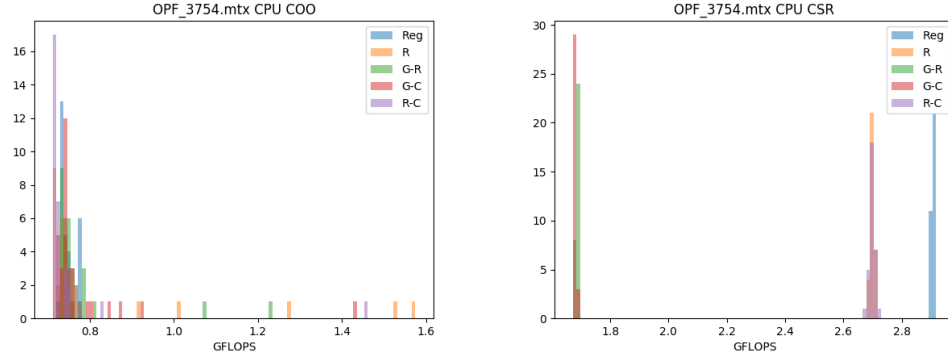


Fig. 6. CPU COO (left) and CPU CSR (left) for OPF 3754

If we take the original matrix and split into parts having the same number of rows, and execute them in parallel using different cores, we can see in Figure 7 that randomization is quite useful.

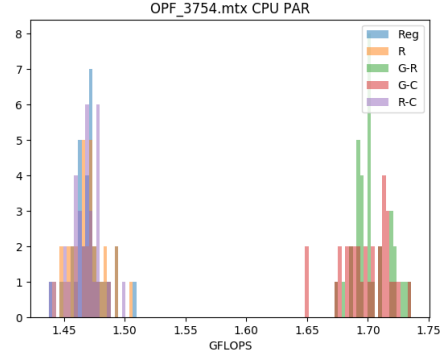


Fig. 7. Parallel CPU CSR for OPF 3754

216 In Figure 8, 9 and 10, randomization is harmful to the GPU implementation. The OPF 375 matrix is mostly diagonal,
 217 thus the vector \mathbf{x} is read in close quarters, randomization breaks it. If the load balance is fixed (i.e., by dividing the
 218 matrix by row and in equal row), randomization is beneficial.

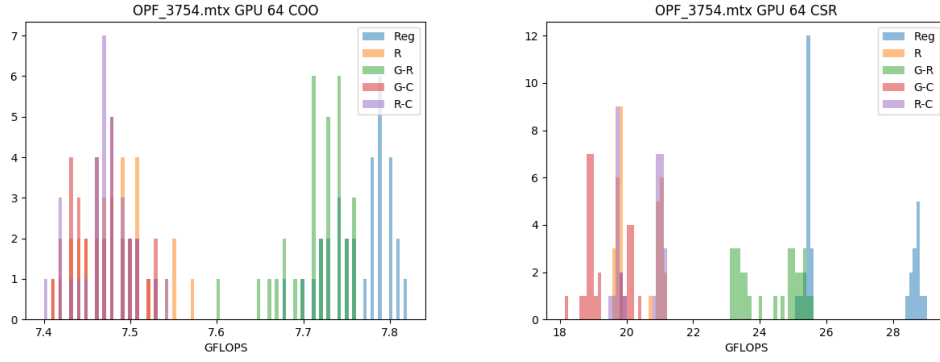


Fig. 8. Vega 20, GPU 64bits COO (left) and GPU CSR (right) for OPF 3754

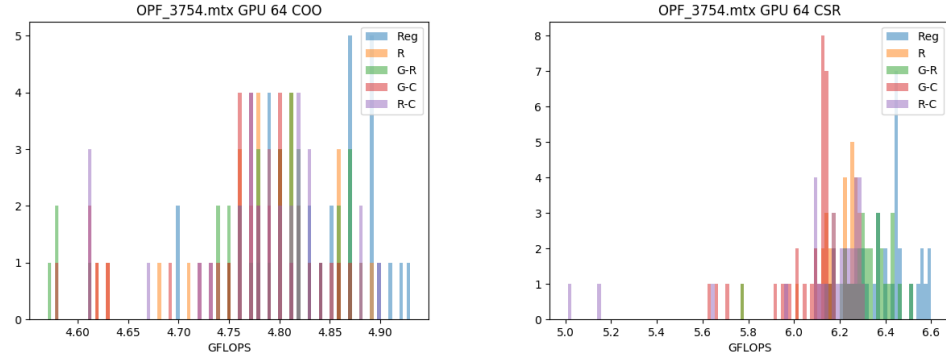


Fig. 9. Ellesmere, GPU 64bits COO (left) and GPU CSR (right) for OPF 3754

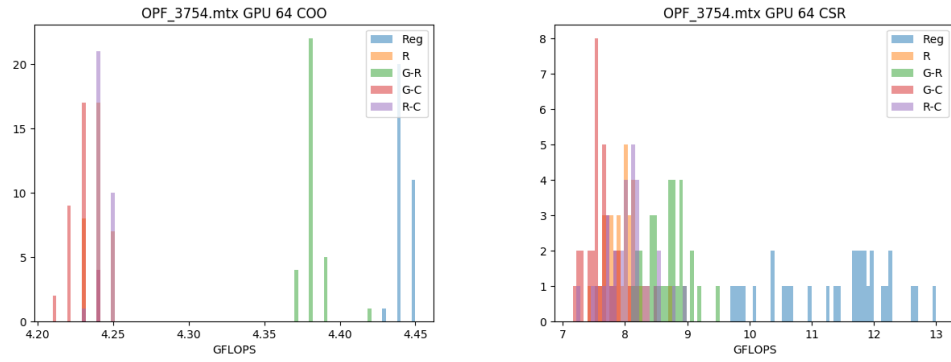


Fig. 10. Fiji, GPU 64bits COO (left) and GPU CSR (right) for OPF 3754

219 For matrix LP OSA 07, randomization helps clearly only for CPU CSR as we show in Figure 11. In Figure 12, 13, and
 220 14, we can see that randomization is harmful but for one GPU, we can show that a single exception is possible (40%
 221 improvement).

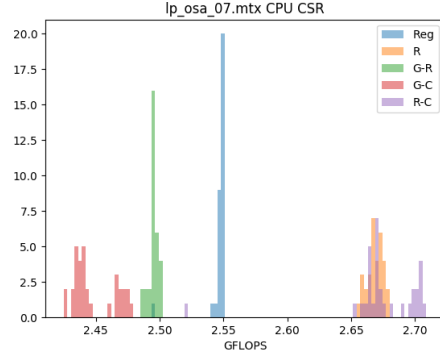


Fig. 11. CPU CSR for LP OSA 07

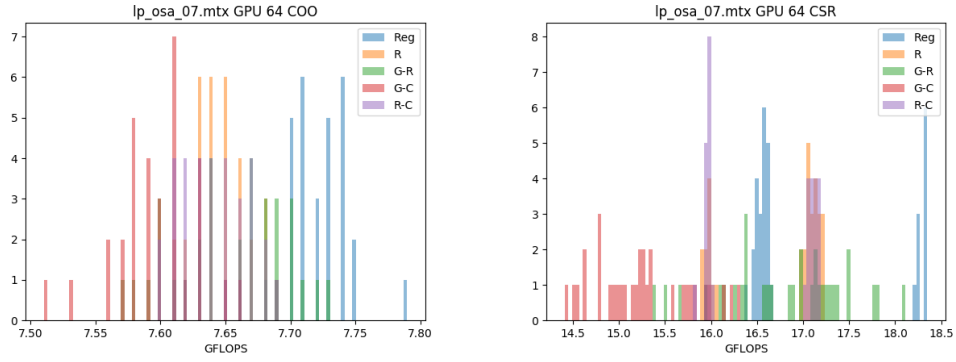


Fig. 12. Vega 20, GPU 64bits COO (left) and GPU CSR (right) for OPF 3754

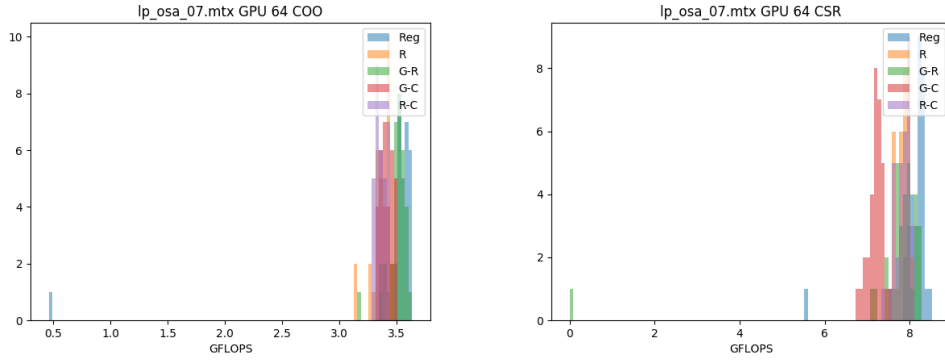


Fig. 13. Ellesmere, GPU 64bits COO (left) and GPU CSR (right) for OPF 3754

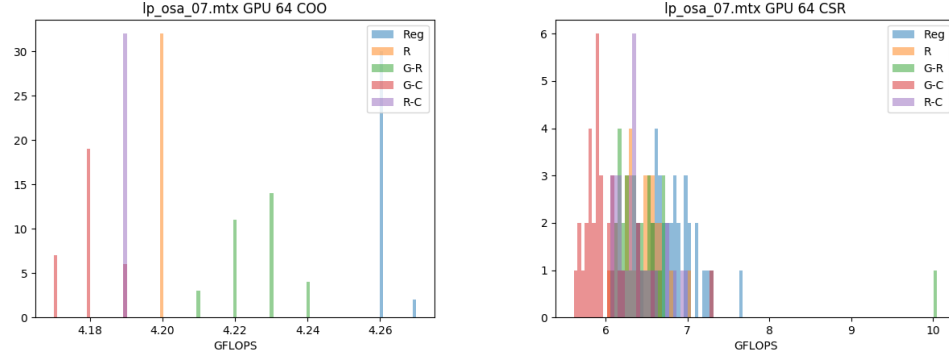


Fig. 14. Fiji, GPU 64bits COO (left) and GPU CSR (right) for OPF 3754

222 An example, the matrix MULT DCOP 01, is where randomization is useful for the CPU, GPU, and the parallel version
 223 Figure 15, 16 - 19 and the gains can be up to 10-15%. Consider, we can achieve these improvements without any insights
 224 to the architecture, the algorithms and their relationships.

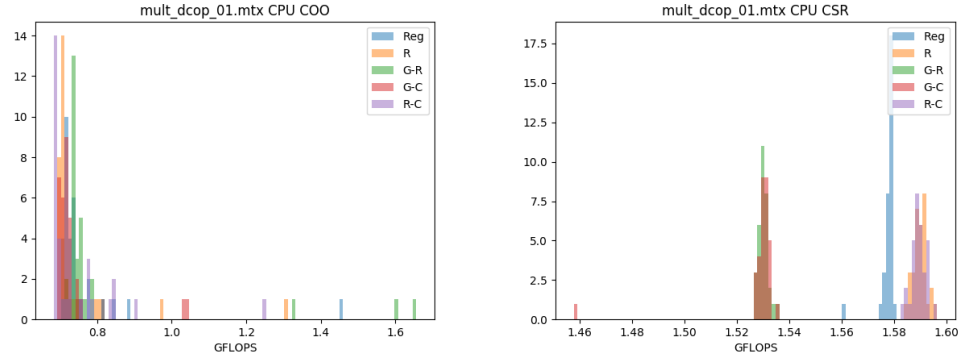


Fig. 15. CPU COO (left) and CPU CSR (right) for MULT DCOP 01

225 What does it mean when randomization does not work? The matrices we use in this work are not chosen randomly
 226 (pun not intended), they are the matrices that are difficult to handle in our custom SpMV engines using a combination
 227 of sorting networks and systolic arrays. If randomization does not work in our simplified work bench, will not work in
 228 our specialized architecture because the reorganization of the matrix or the input and output vector does not have
 229 the necessary parallelism, data locality, and data streaming. We need to do something else. In this case disrupting the
 230 memory pattern is not sufficient. Thus, if we cannot beat the pattern, we must exploit it, well not in this work.

231 7 WORKLOADS

232 In the previous sections, we defined what we mean for randomization and we present our tools of tricks for the measure
 233 of the effects of randomization. Here we describe the work loads, the applications, we use to test the effects of the
 234 randomization.

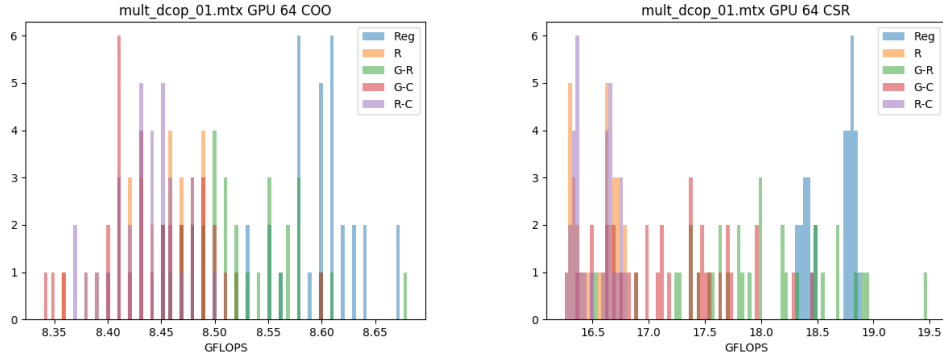


Fig. 16. Vega 20, GPU 64bits COO (left) and GPU CSR (right) for MULT DCOP 01

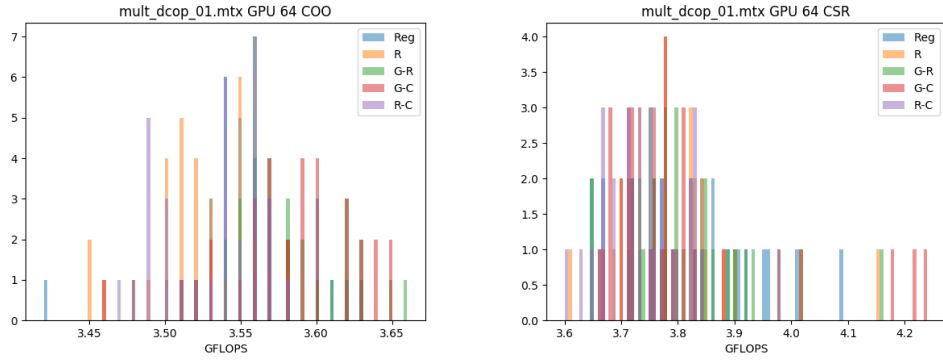


Fig. 17. Ellesmere, GPU 64bits COO (left) and GPU CSR (right) for MULT DCOP 01

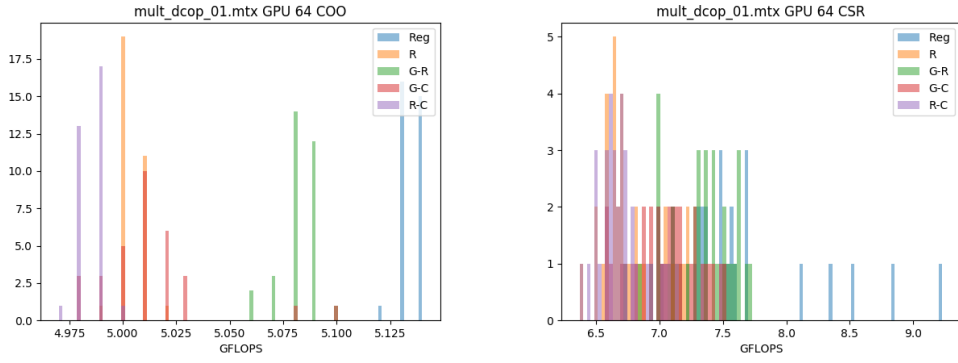


Fig. 18. Fiji, GPU 64bits COO (left) and GPU CSR (right) for MULT DCOP 01

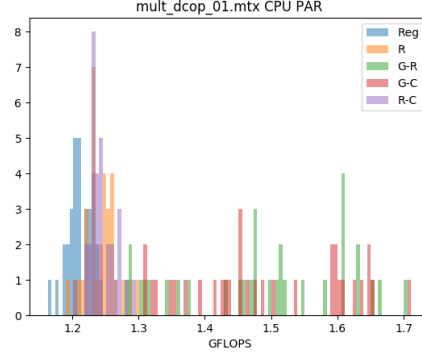


Fig. 19. Parallel CPU CSR for MULT DCOP 01

7.1 Python COO and CSR algorithms

The simplicity to compute the SpMV by the code $z = A * b$ in Python is very rewarding. By change of the matrix storage format, $A = A.tocsr(); z = A * b$, we have a different algorithm. The performance exploitation is moved to the lower level. The CSR implementation is often two times faster but there are edge cases where the COO and COO with randomization can go beyond and be surprisingly better: MUL DCOP 03 is an example where COO can do well.

Intuitively, Randomization can affect the performance because the basic implementation is a sorting algorithm and it is a fixed algorithm. There are many sorting algorithms and each can be optimal for a different initial distribution. If we knew what is the sorting algorithm we could tailor the input distribution. Here we just play with it.

In Section 8, we present all the results for CPU and GPUS. Keep in mind that these problems are hard, in the sense they do not have fancy performance sheets (these architectures can achieve Tera FLOPs sustained performance for dense computations). If we go through diligently, we can see that there is a 15x performance difference between the single thread CPU and Vega 20 GPU (i.e, 3 vs 40 GFLOPS).

7.2 Parallel CSR using up to 16 cores

Python provides the concept of Pool to exploit a naive parallel computation. We notice that work given to a Pool is split accordingly to the number of elements to separate HW cores. We also noticed that the work load move from a core to another, thus not ideal. Also we notice that Pool introduce a noticeable overhead: a Pool of 1, never achieves the performance of the single thread $z = A * b$. Using Pool allows us to investigate how a naive row partitioning without counting can scale up with number of cores. We tested by splitting the rows to 1–16 cores evenly (one thread per core) and we present the performance for only the best configuration. The randomization goal is to distribute the work uniformly: a balanced work distribution avoid the unfortunate case where a single core does all the work. We are pleased by the simplicity of the benchmark and we know we can do better.

7.3 GPU COO and CSR algorithms

In this work, we use AMD GPUs and *rocSPARSE* is their current software. The software has a few glitches but overall can be used for different generation of AMD GPUs. We use the COO and CSR algorithms and we provide performance measure for double precision only. The ideas of using different GPUs: it is important to verify that the randomization

can be applied independently of the HW. We are not here to compare performance across GPUs and CPUs. Often the limitation is the software, how the software can exploit the hardware or how the software will make easy to use a specific GPU. For example, the Fiji architecture is clearly superior to the Ellesmere, however the latter have better support and the system overall is more stable and user friendly.

The performance of the CSR algorithm is about two times faster than the COO. Most of the algorithms count the number of sparse elements in a row and thus they can decide the work load partition accordingly. Counting give you an edge but without changing the order of the computation there could be cases where the work load is not balanced and a little randomization could help and it does.

7.4 Randomization sometimes works

For the majority of the cases we investigated and reported in the following sections, Randomization does not work. However, there are cases where randomization does work and does work for different algorithms and architectures. If you are in the business of preconditioning, permutations are pretty cheap. If you can find a good one just consider like a preconditioning matrix, which it is.

This shows also that HW has to be more conscious, well the HW designer should, and accept that there are options at software level, at matrix level and beyond.

8 EXPERIMENTAL RESULTS

The main hardware setup is a AMD Threadripper with 16 cores. We have three Radeon GPUs: Vega 20 7nm, Pro 2xFiji, and Pro 2xEllesmere.

Vega 20 can deliver 3.5TFLOPS in double precision and it has 1TB/s HBM memory. Each Fiji provides 0.5 TFLOPS in double precision and has 512GB/s HBM, the card has two chips. The Ellesmere provides 0.3TFLOPS in double precision and has 224GB/s DDR5, the card has two chips. In the performance plots presented earlier and in the following, you will notice that the performance gap between these GPUs is not so marked. We can safely state that $vega \sim 2 \times Fiji$ and $Fiji \sim 2 \times ellesmere$

There are 4 basic randomization formats:

- **Random Row Permutation**, we take the original matrix and permute the rows.
- **Random Row and Column Permutation**, we take the original matrix and permute the rows and the columns.
- **Gradient based row permutation**, we compute the row histogram and we compute the gradient: $h_{i+1} - h_i$. We find a single point where the gradient is maximum, this is the pivot for a shuffle like a magician would shuffle a deck of cards. Then we permute the two parts randomly.
- **Gradient based row and column permutation**, As above but also for the columns.

For large matrices (large number of columns and rows) a permutation tends to be a close variation of the original, still a random permutation. The gradient allows us to describe two areas of the original matrix where there is a clear and de-marked density variation: for example, there are two uniform distributed sub matrices but one denser than the other. A shuffle redistributes every other sample/card to different parts and these can be permuted locally.

We report in the following the performance results GFLOPS, we introduce a * following the best performance. This is tedious to read and, we assure, to write. The code and the results are available as software repository. Remember each experiment is based on 32 different runs and thus we report maximum, minimum, and mean as a summary. We use the symbol H for entropy.

9 VEGA VII AND THREADRIPPER

298	mult_dcop_03.mtx
299	Regular
300	
301	CPU COO min 0.728 max 0.880 mean 0.757
302	CPU CSR min 1.563 max 1.581 mean 1.577
303	GPU 64 COO min 8.540 max* 8.670 mean 8.619
304	CSR min 18.320 max 18.930 mean 18.620
305	CPU PAR min 1.170 max 1.269 mean 1.226
306	H min 9.689 max 9.689 mean 9.689
307	Row-Premute
308	CPU COO min 0.710 max 0.845 mean 0.724
309	CPU CSR min 1.549 max* 1.597 mean 1.589
310	GPU 64 COO min 8.360 max 8.540 mean 8.442
311	CSR min 16.260 max 16.780 mean 16.551
312	CPU PAR min 1.205 max 1.319 mean 1.263
313	H min 10.737 max 10.742 mean 10.740
314	Row-Gradient
315	CPU COO min 0.706 max 1.603 mean 0.806
316	CPU CSR min 1.493 max 1.534 mean 1.528
317	GPU 64 COO min 8.430 max 8.610 mean 8.527
318	CSR min 17.070 max*18.970 mean 18.115
319	CPU PAR min 1.331 max 1.695 mean 1.513
320	H min 10.576 max 10.585 mean 10.580
321	Column-Gradient
322	CPU COO min 0.694 max* 1.632 mean 0.797
323	CPU CSR min 1.491 max 1.534 mean 1.529
324	GPU 64 COO min 8.350 max 8.520 mean 8.429
325	CSR min 15.970 max 18.180 mean 17.124
326	CPU PAR min 1.321 max* 1.728 mean 1.514
327	H min 10.826 max*10.840 mean 10.833
328	Row-Column-Permute
329	CPU COO min 0.688 max 0.757 mean 0.696
330	CPU CSR min 1.490 max 1.595 mean 1.584
331	GPU 64 COO min 8.380 max 8.500 mean 8.445
332	CSR min 16.230 max 16.780 mean 16.513
333	CPU PAR min 1.192 max 1.274 mean 1.237
334	H min 10.737 max 10.742 mean 10.740
335	mult_dcop_01.mtx
336	Regular
337	CPU COO min 0.710 max 1.453 mean 0.761
338	CPU CSR min 1.561 max 1.581 mean 1.578
339	GPU 64 COO min 8.520 max 8.670 mean 8.597
340	CSR min 18.320 max 18.870 mean 18.636
341	CPU PAR min 1.163 max 1.246 mean 1.212
342	H min 9.689 max 9.689 mean 9.689
343	Row-Premute
344	CPU COO min 0.699 max 1.305 mean 0.745
345	CPU CSR min 1.585 max 1.597 mean 1.590
346	GPU 64 COO min 8.360 max 8.520 mean 8.446
347	CSR min 16.260 max 16.780 mean 16.528
348	CPU PAR min 1.192 max 1.298 mean 1.242
349	H min 10.738 max 10.742 mean 10.740
350	Row-Gradient
351	CPU COO min 0.709 max* 1.656 mean 0.819
352	CPU CSR min 1.527 max 1.535 mean 1.530
353	GPU 64 COO min 8.450 max* 8.680 mean 8.527
354	CSR min 16.520 max*19.480 mean 17.984
355	CPU PAR min 1.280 max 1.704 mean 1.485
356	H min 10.572 max 10.585 mean 10.581
357	Column-Gradient
358	CPU COO min 0.698 max 1.042 mean 0.737
359	CPU CSR min 1.458 max 1.536 mean 1.528
360	GPU 64 COO min 8.340 max 8.600 mean 8.443
361	CSR min 16.360 max 18.450 mean 17.247
362	CPU PAR min 1.307 max* 1.712 mean 1.494
363	H min 10.823 max*10.841 mean 10.835
364	Row-Column-Permute
365	CPU COO min 0.683 max 1.247 mean 0.749
366	CPU CSR min 1.583 max* 1.595 mean 1.590
367	GPU 64 COO min 8.370 max 8.500 mean 8.435
368	CSR min 16.250 max 16.780 mean 16.518
369	CPU PAR min 1.206 max 1.291 mean 1.243
370	H min 10.738 max 10.742 mean 10.740

371	mult_dcop_02.mtx
372	Regular
373	
374	CPU COO min 1.615 max* 1.677 mean 1.652
375	CPU CSR min 1.539 max 1.579 mean 1.575
376	GPU 64 COO min 8.530 max* 8.700 mean 8.614
377	CSR min 18.290 max 18.890 mean 18.597
378	CPU PAR min 1.120 max 1.248 mean 1.211
379	H min 9.689 max 9.689 mean 9.689
380	Row-Premute
381	CPU COO min 0.684 max 0.780 mean 0.705
382	CPU CSR min 1.558 max* 1.596 mean 1.588
383	GPU 64 COO min 8.360 max 8.490 mean 8.433
384	CSR min 16.240 max 16.750 mean 16.552
385	CPU PAR min 1.182 max 1.277 mean 1.242
386	H min 10.737 max 10.742 mean 10.740
387	Row-Gradient
388	CPU COO min 0.704 max 1.373 mean 0.790
389	CPU CSR min 1.518 max 1.535 mean 1.529
390	GPU 64 COO min 8.420 max 8.590 mean 8.517
391	CSR min 16.680 max*19.550 mean 17.907
392	CPU PAR min 1.328 max* 1.713 mean 1.484
393	H min 10.572 max 10.585 mean 10.581
394	Column-Gradient
395	CPU COO min 0.697 max 1.460 mean 0.742
396	CPU CSR min 1.517 max 1.534 mean 1.527
397	GPU 64 COO min 8.330 max 8.490 mean 8.420
398	CSR min 16.020 max 18.390 mean 17.303
399	CPU PAR min 1.321 max 1.709 mean 1.557
400	H min 10.823 max*10.843 mean 10.835
401	Row-Column-Permute
402	CPU COO min 0.691 max 0.746 mean 0.698
403	CPU CSR min 1.568 max 1.595 mean 1.587
404	GPU 64 COO min 8.350 max 8.500 mean 8.436
405	CSR min 16.250 max 16.780 mean 16.517
406	CPU PAR min 1.187 max 1.280 mean 1.228
407	H min 10.739 max 10.743 mean 10.740
408	lp_fit2d.mtx
409	Regular
410	CPU COO min 0.774 max 0.804 mean 0.793
411	CPU CSR min 2.538 max 2.550 mean 2.547
412	GPU 64 COO min 7.060 max 7.170 mean 7.101
413	CSR min 15.650 max*18.700 mean 18.031
414	CPU PAR min 1.537 max 1.645 mean 1.590
415	H min 11.109 max 11.109 mean 11.109
416	Row-Premute
417	CPU COO min 0.740 max 0.776 mean 0.746
418	CPU CSR min 3.302 max* 3.328 mean 3.317
419	GPU 64 COO min 7.040 max* 7.180 mean 7.098
420	CSR min 15.690 max 18.580 mean 16.732
421	CPU PAR min 1.327 max 1.482 mean 1.422
422	H min 11.098 max 11.105 mean 11.101
423	Row-Gradient
424	CPU COO min 0.739 max* 2.092 mean 1.091
425	CPU CSR min 2.539 max 2.546 mean 2.543
426	GPU 64 COO min 7.040 max 7.150 mean 7.100
427	CSR min 15.520 max 18.560 mean 17.547
428	CPU PAR min 1.401 max 1.661 mean 1.525
429	H min 11.109 max 11.109 mean 11.109
430	Column-Gradient
431	CPU COO min 0.726 max 2.065 mean 1.011
432	CPU CSR min 2.539 max 2.550 mean 2.546
433	GPU 64 COO min 6.800 max 7.140 mean 7.080
434	CSR min 15.480 max 18.560 mean 16.866
435	CPU PAR min 1.391 max* 1.737 mean 1.563
436	H min 11.329 max 11.333 mean 11.331
437	Row-Column-Permute
438	CPU COO min 0.746 max 0.782 mean 0.754
439	CPU CSR min 3.310 max 3.324 mean 3.318
440	GPU 64 COO min 7.030 max 7.160 mean 7.100
441	CSR min 15.730 max 18.530 mean 17.362
442	CPU PAR min 1.340 max 1.451 mean 1.401
443	H min 11.099 max 11.104 mean 11.102
444	bloweya.mtx
	Regular

445		CPU COO	min	0.727	max*	1.815	mean	0.892	519		GPU 64 COO	min	11.340	max*	11.860	mean	11.441
446		CPU CSR	min	2.867	max*	2.936	mean	2.917	520		CSR	min	36.010	max*	40.960	mean	38.048
447		GPU 64 COO	min	0.000	max	0.000	mean	0.000	521		CPU PAR	min	2.019	max	2.204	mean	2.130
448		CSR	min	0.000	max	0.000	mean	0.000	522		H	min	8.228	max	8.228	mean	8.228
449		CPU PAR	min	1.680	max*	1.751	mean	1.719	523	Row-Premute							
450		H	min	7.205	max	7.205	mean	7.205	524		CPU COO	min	0.718	max	0.751	mean	0.732
451	Row-Premute								525		CPU CSR	min	2.488	max	2.507	mean	2.498
452		CPU COO	min	0.678	max	1.483	mean	0.746	526		GPU 64 COO	min	10.810	max	11.090	mean	10.949
453		CPU CSR	min	2.311	max	2.326	mean	2.320	527		CSR	min	24.860	max	26.410	mean	25.527
454		GPU 64 COO	min	6.840	max*	7.270	mean	6.930	528		CPU PAR	min	1.978	max	2.290	mean	2.135
455		CSR	min	15.650	max	16.800	mean	16.233	529		H	min	11.836	max	11.840	mean	11.838
456		CPU PAR	min	1.649	max	1.730	mean	1.682	530	Row-Gradient							
457		H	min	11.026	max	11.031	mean	11.029	531		CPU COO	min	0.722	max	1.794	mean	0.769
458	Row-Gradient								532		CPU CSR	min	2.407	max	2.421	mean	2.416
459		CPU COO	min	0.708	max	1.209	mean	0.779	533		GPU 64 COO	min	11.210	max	11.480	mean	11.317
460		CPU CSR	min	1.648	max	1.735	mean	1.709	534		CSR	min	31.920	max	34.690	mean	33.246
461		GPU 64 COO	min	6.920	max	7.080	mean	7.015	535		CPU PAR	min	2.184	max*	2.302	mean	2.232
462		CSR	min	16.950	max	19.500	mean	17.794	536		H	min	10.742	max	10.757	mean	10.748
463		CPU PAR	min	1.497	max	1.743	mean	1.608	537	Column-Gradient							
464		H	min	10.298	max	10.304	mean	10.301	538		CPU COO	min	0.720	max	0.916	mean	0.742
465	Column-Gradient								539		CPU CSR	min	2.395	max	2.410	mean	2.402
466		CPU COO	min	0.709	max	1.536	mean	0.817	540		GPU 64 COO	min	10.840	max	11.070	mean	10.946
467		CPU CSR	min	1.705	max	1.753	mean	1.735	541		CSR	min	24.340	max	26.140	mean	25.393
468		GPU 64 COO	min	6.800	max	7.120	mean	6.865	542		CPU PAR	min	2.184	max	2.272	mean	2.223
469		CSR	min	15.480	max*	17.710	mean	16.470	543		H	min	11.873	max	11.882	mean	11.878
470		CPU PAR	min	1.446	max	1.718	mean	1.591	544	Row-Column-Permute							
471		H	min	10.880	max	10.886	mean	10.883	545		CPU COO	min	0.707	max	0.748	mean	0.714
472	Row-Column-Permute								546		CPU CSR	min	2.458	max	2.511	mean	2.506
473		CPU COO	min	0.670	max	1.024	mean	0.706	547		GPU 64 COO	min	10.880	max	11.070	mean	10.957
474		CPU CSR	min	2.199	max	2.340	mean	2.326	548		CSR	min	24.890	max	26.490	mean	25.642
475		GPU 64 COO	min	6.800	max	6.980	mean	6.933	549		CPU PAR	min	2.209	max	2.282	mean	2.240
476		CSR	min	15.610	max	16.900	mean	16.227	550		H	min	11.834	max*	11.840	mean	11.838
477		CPU PAR	min	1.598	max	1.668	mean	1.632	551	brainpc2.mtx							
478		H	min	11.025	max*	11.032	mean	11.029	552	Regular							
479	lp_osa_07.mtx								553		CPU COO	min	0.732	max	0.751	mean	0.744
480	Regular								554		CPU CSR	min	2.885	max*	2.916	mean	2.909
481		CPU COO	min	0.715	max	1.798	mean	0.885	555		GPU 64 COO	min	0.000	max	0.000	mean	0.000
482		CPU CSR	min	2.495	max	2.551	mean	2.547	556		CSR	min	0.000	max	0.000	mean	0.000
483		GPU 64 COO	min	7.650	max*	7.790	mean	7.718	557		CPU PAR	min	1.276	max	1.299	mean	1.286
484		CSR	min	16.390	max*	18.350	mean	17.093	558		H	min	7.478	max	7.478	mean	7.478
485		CPU PAR	min	0.963	max	1.012	mean	0.995	559	Row-Premute							
486		H	min	8.412	max	8.412	mean	8.412	560		CPU COO	min	0.727	max	0.855	mean	0.736
487	Row-Premute								561		CPU CSR	min	2.385	max	2.411	mean	2.397
488		CPU COO	min	0.720	max*	2.078	mean	1.104	562		GPU 64 COO	min	8.120	max	8.410	mean	8.206
489		CPU CSR	min	2.656	max*	2.679	mean	2.669	563		CSR	min	18.670	max	19.960	mean	19.536
490		GPU 64 COO	min	7.610	max	7.690	mean	7.647	564		CPU PAR	min	1.293	max	1.340	mean	1.314
491		CSR	min	15.910	max	17.210	mean	16.750	565		H	min	9.809	max	9.813	mean	9.811
492		CPU PAR	min	0.890	max	0.940	mean	0.918	566	Row-Gradient							
493		H	min	9.255	max	9.258	mean	9.256	567		CPU COO	min	0.696	max*	1.546	mean	0.785
494	Row-Gradient								568		CPU CSR	min	1.361	max	1.420	mean	1.411
495		CPU COO	min	0.725	max	2.078	mean	1.041	569		GPU 64 COO	min	8.190	max*	8.550	mean	8.302
496		CPU CSR	min	2.487	max	2.502	mean	2.495	570		CSR	min	18.700	max*	21.000	mean	19.890
497		GPU 64 COO	min	7.570	max	7.730	mean	7.655	571		CPU PAR	min	1.435	max	1.666	mean	1.549
498		CSR	min	15.370	max	18.100	mean	16.803	572		H	min	9.721	max	9.727	mean	9.723
499		CPU PAR	min	1.435	max	1.796	mean	1.592	573	Column-Gradient							
500		H	min	8.637	max	8.678	mean	8.672	574		CPU COO	min	0.698	max	1.467	mean	0.746
501	Column-Gradient								575		CPU CSR	min	1.377	max	1.423	mean	1.414
502		CPU COO	min	0.724	max	1.990	mean	1.000	576		GPU 64 COO	min	8.110	max	8.290	mean	8.187
503		CPU CSR	min	2.425	max	2.477	mean	2.448	577		CSR	min	18.090	max	20.190	mean	19.217
504		GPU 64 COO	min	7.510	max	7.660	mean	7.596	578		CPU PAR	min	1.345	max*	1.681	mean	1.518
505		CSR	min	14.410	max	16.290	mean	15.267	579		H	min	10.369	max*	10.372	mean	10.370
506		CPU PAR	min	1.238	max	1.774	mean	1.534	580	Row-Column-Permute							
507		H	min	9.447	max*	9.603	mean	9.576	581		CPU COO	min	0.698	max	1.390	mean	0.788
508	Row-Column-Permute								582		CPU CSR	min	2.387	max	2.410	mean	2.399
509		CPU COO	min	0.738	max	1.950	mean	1.071	583		GPU 64 COO	min	8.120	max	8.260	mean	8.191
510		CPU CSR	min	2.522	max	2.709	mean	2.675	584		CSR	min	18.530	max	19.960	mean	19.307
511		GPU 64 COO	min	7.600	max	7.690	mean	7.641	585		CPU PAR	min	1.295	max	1.347	mean	1.319
512		CSR	min	15.820	max	17.190	mean	16.572	586		H	min	9.809	max	9.813	mean	9.811
513		CPU PAR	min	0.891	max	0.944	mean	0.924	587	shermanACb.mtx							
514		H	min	9.255	max	9.258	mean	9.256	588	Regular							
515	ex19.mtx								589		CPU COO	min	0.712	max	1.201	mean	0.756
516	Regular								590		CPU CSR	min	1.558	max	1.601	mean	1.596
517		CPU COO	min	0.732	max*	1.837	mean	1.076	591		GPU 64 COO	min	7.080	max*	7.370	mean	7.184
518		CPU CSR	min	2.563	max*	2.586	mean	2.577	592		CSR	min	17.580	max*	19.480	mean	18.770

		CPU PAR	min	1.286	max	1.511	mean	1.447	667	Row-Permute										
594		H	min	8.600	max	8.600	mean	8.600	668		CPU COO	min	0.724	max	1.100	mean	0.765			
595	Row-Permute								669		CPU CSR	min	2.581	max	2.626	mean	2.609			
596		CPU COO	min	0.689	max	0.890	mean	0.704	670		GPU 64 COO	min	7.170	max	7.340	mean	7.253			
597		CPU CSR	min	1.600	max	1.630	mean	1.618	671		CSR min	17.360	max	18.500	mean	18.014				
598		GPU 64 COO	min	7.000	max	7.180	mean	7.061	672		CPU PAR	min	1.494	max	1.607	mean	1.558			
599		CSR	min	15.760	max	17.240	mean	16.625	673		H	min	10.043	max	10.047	mean	10.044			
600		CPU PAR	min	1.296	max	1.419	mean	1.365	674	Row-Gradient										
601		H	min	10.376	max	10.380	mean	10.379	675		CPU COO	min	0.716	max	1.701	mean	0.804			
602	Row-Gradient								676		CPU CSR	min	1.824	max	1.840	mean	1.832			
603		CPU COO	min	0.704	max	1.615	mean	0.806	677		GPU 64 COO	min	7.228	max	7.510	mean	7.383			
604		CPU CSR	min	1.355	max	1.370	mean	1.362	678		CSR min	17.540	max	20.710	mean	19.302				
605		GPU 64 COO	min	7.020	max	7.160	mean	7.083	679		CPU PAR	min	1.384	max	1.593	mean	1.526			
606		CSR	min	0.000	max	16.290	mean	15.076	680		H	min	9.681	max	9.706	mean	9.694			
607		CPU PAR	min	1.256	max	1.520	mean	1.405	681	Column-Gradient										
608		H	min	9.915	max	9.925	mean	9.921	682		CPU COO	min	0.711	max	1.029	mean	0.746			
609	Column-Gradient								683		CPU CSR	min	1.817	max	1.834	mean	1.827			
610		CPU COO	min	0.702	max	1.626	mean	0.844	684		GPU 64 COO	min	7.110	max	7.270	mean	7.193			
611		CPU CSR	min	1.327	max	1.374	mean	1.364	685		CSR min	16.530	max	18.590	mean	17.574				
612		GPU 64 COO	min	6.920	max	7.210	mean	7.030	686		CPU PAR	min	1.390	max	1.574	mean	1.511			
613		CSR	min	0.000	max	15.260	mean	14.279	687		H	min	10.612	max	10.659	mean	10.634			
614		CPU PAR	min	1.283	max	1.531	mean	1.385	688	Row-Column-Permute										
615		H	min	10.572	max	10.595	mean	10.590	689		CPU COO	min	0.719	max	1.391	mean	0.756			
616	Row-Column-Permute								690		CPU CSR	min	2.546	max	2.625	mean	2.611			
617		CPU COO	min	0.707	max	1.532	mean	0.924	691		GPU 64 COO	min	7.190	max	7.320	mean	7.248			
618		CPU CSR	min	1.606	max	1.634	mean	1.624	692		CSR min	17.500	max	18.640	mean	18.040				
619		GPU 64 COO	min	6.970	max	7.110	mean	7.045	693		CPU PAR	min	1.465	max	1.573	mean	1.533			
620		CSR	min	15.850	max	17.310	mean	16.783	694		H	min	10.041	max	10.046	mean	10.044			
621		CPU PAR	min	1.286	max	1.406	mean	1.357	695	TSOPFF_FS_b9_c6.mtx										
622		H	min	10.377	max	10.382	mean	10.379	696	Regular										
623	cvxqp3.mtx								697		CPU COO	min	0.705	max	0.734	mean	0.718			
624	Regular								698		CPU CSR	min	3.028	max	3.052	mean	3.045			
625		CPU COO	min	0.697	max	0.720	mean	0.712	699		GPU 64 COO	min	0.000	max	0.000	mean	0.000			
626		CPU CSR	min	2.624	max	2.643	mean	2.638	700		CSR min	0.000	max	0.000	mean	0.000				
627		GPU 64 COO	min	6.060	max	6.220	mean	6.121	701		CPU PAR	min	1.528	max	1.602	mean	1.568			
628		CSR	min	19.450	max	22.710	mean	21.277	702		H	min	7.380	max	7.380	mean	7.380			
629		CPU PAR	min	1.733	max	1.860	mean	1.804	703	Row-Permute										
630		H	min	8.646	max	8.646	mean	8.646	704		CPU COO	min	0.733	max	1.640	mean	0.777			
631	Row-Permute								705		CPU CSR	min	2.450	max	2.543	mean	2.525			
632		CPU COO	min	0.695	max	1.577	mean	0.894	706		GPU 64 COO	min	7.200	max	7.320	mean	7.268			
633		CPU CSR	min	2.452	max	2.471	mean	2.464	707		CSR min	17.420	max	18.540	mean	18.102				
634		GPU 64 COO	min	5.870	max	6.060	mean	5.930	708		CPU PAR	min	1.474	max	1.595	mean	1.546			
635		CSR	min	17.510	max	19.130	mean	18.516	709		H	min	10.042	max	10.046	mean	10.044			
636		CPU PAR	min	1.723	max	1.833	mean	1.774	710	Row-Gradient										
637		H	min	11.028	max	11.033	mean	11.030	711		CPU COO	min	0.712	max	0.926	mean	0.750			
638	Row-Gradient								712		CPU CSR	min	1.819	max	1.846	mean	1.832			
639		CPU COO	min	0.693	max	1.523	mean	0.788	713		GPU 64 COO	min	7.210	max	7.370	mean	7.298			
640		CPU CSR	min	1.287	max	1.305	mean	1.296	714		CSR min	17.550	max	20.740	mean	19.089				
641		GPU 64 COO	min	5.920	max	6.000	mean	5.962	715		CPU PAR	min	1.256	max	1.554	mean	1.495			
642		CSR	min	16.810	max	18.410	mean	17.561	716		H	min	9.666	max	9.704	mean	9.690			
643		CPU PAR	min	1.378	max	1.485	mean	1.429	717	Column-Gradient										
644		H	min	11.061	max	11.069	mean	11.064	718		CPU COO	min	0.710	max	1.690	mean	0.791			
645	Column-Gradient								719		CPU CSR	min	1.813	max	1.836	mean	1.830			
646		CPU COO	min	0.693	max	1.521	mean	0.772	720		GPU 64 COO	min	7.130	max	7.310	mean	7.211			
647		CPU CSR	min	1.291	max	1.302	mean	1.297	721		CSR min	16.550	max	18.690	mean	17.617				
648		GPU 64 COO	min	5.900	max	6.060	mean	5.960	722		CPU PAR	min	1.385	max	1.539	mean	1.506			
649		CSR	min	16.620	max	18.330	mean	17.592	723		H	min	10.611	max	10.659	mean	10.634			
650		CPU PAR	min	1.372	max	1.464	mean	1.409	724	Row-Column-Permute										
651		H	min	11.127	max	11.135	mean	11.130	725		CPU COO	min	0.709	max	1.531	mean	0.963			
652	Row-Column-Permute								726		CPU CSR	min	2.506	max	2.648	mean	2.622			
653		CPU COO	min	0.704	max	1.503	mean	0.875	727		GPU 64 COO	min	7.140	max	7.330	mean	7.244			
654		CPU CSR	min	2.447	max	2.468	mean	2.459	728		CSR min	17.410	max	18.520	mean	18.148				
655		GPU 64 COO	min	5.880	max	5.980	mean	5.931	729		CPU PAR	min	1.466	max	1.574	mean	1.528			
656		CSR	min	17.550	max	19.140	mean	18.227	730		H	min	10.041	max	10.046	mean	10.044			
657		CPU PAR	min	1.639	max	1.743	mean	1.704	731	OPF_6000.mtx										
658		H	min	11.028	max	11.035	mean	11.030	732	Regular										
659	case9.mtx								733		CPU COO	min	0.714	max	0.731	mean	0.720			
660	Regular								734		CPU CSR	min	2.667	max	2.770	mean	2.720			
661		CPU COO	min	0.721	max	1.800	mean	1.177	735		GPU 64 COO	min	12.310	max	12.550	mean	12.425			
662		CPU CSR	min	3.021	max	3.046	mean	3.036	736		CSR min	39.860	max	43.770	mean	42.075				
663		GPU 64 COO	min	0.000	max	0.000	mean	0.000	737		CPU PAR	min	1.735	max	1.945	mean	1.845			
664		CSR	min	0.000	max	0.000	mean	0.000	738		H	min	8.799	max	8.799	mean	8.799			
665		CPU PAR	min	1.508	max	1.605	mean	1.573	739	Row-Permute										
666		H	min	7.380	max	7.380	mean	7.380	740		CPU COO	min	0.689	max	0.710	mean	0.695			

741		CPU CSR	min	2.358	max	2.413	mean	2.392	815		CSR	min	19.960	max	21.190	mean	20.696
742		GPU 64 COO	min	11.430	max	11.770	mean	11.549	816		CPU PAR	min	1.303	max	1.371	mean	1.345
743		CSR	min	24.470	max	25.580	mean	24.785	817		H	min	10.059	max	10.062	mean	10.061
744		CPU PAR	min	1.758	max	1.896	mean	1.829	818	Row-Gradient							
745		H	min	11.872	max	11.877	mean	11.875	819		CPU COO	min	0.723	max	0.984	mean	0.753
746	Row-Gradient								820		CPU CSR	min	1.781	max	1.809	mean	1.803
747		CPU COO	min	0.716	max	0.775	mean	0.739	821		GPU 64 COO	min	9.380	max	9.660	mean	9.464
748		CPU CSR	min	1.651	max	1.689	mean	1.675	822		CSR	min	15.770	max	19.090	mean	18.037
749		GPU 64 COO	min	12.100	max	12.410	mean	12.205	823		CPU PAR	min	1.775	max	1.924	mean	1.868
750		CSR	min	31.670	max	34.910	mean	33.370	824		H	min	10.205	max	10.233	mean	10.219
751		CPU PAR	min	2.079	max	2.286	mean	2.207	825	Column-Gradient							
752		H	min	11.111	max	11.116	mean	11.113	826		CPU COO	min	0.715	max	0.926	mean	0.757
753	Column-Gradient								827		CPU CSR	min	1.729	max	1.802	mean	1.791
754		CPU COO	min	0.715	max	1.021	mean	0.743	828		GPU 64 COO	min	9.080	max	9.270	mean	9.158
755		CPU CSR	min	1.655	max	1.674	mean	1.666	829		CSR	min	13.980	max	15.780	mean	14.938
756		GPU 64 COO	min	11.340	max	11.560	mean	11.463	830		CPU PAR	min	1.751	max	1.906	mean	1.846
757		CSR	min	23.770	max	25.470	mean	24.489	831		H	min	11.213	max	11.232	mean	11.222
758		CPU PAR	min	2.056	max	2.172	mean	2.118	832	Row-Column-Permute							
759		H	min	12.040	max	12.047	mean	12.043	833		CPU COO	min	0.732	max	1.598	mean	0.785
760	Row-Column-Permute								834		CPU CSR	min	2.594	max	2.602	mean	2.599
761		CPU COO	min	0.677	max	0.785	mean	0.687	835		GPU 64 COO	min	9.340	max	9.460	mean	9.394
762		CPU CSR	min	2.325	max	2.434	mean	2.369	836		CSR	min	19.950	max	21.500	mean	20.544
763		GPU 64 COO	min	11.450	max	11.650	mean	11.538	837		CPU PAR	min	1.326	max	1.374	mean	1.354
764		CSR	min	24.330	max	25.560	mean	25.008	838		H	min	10.059	max	10.062	mean	10.061
765		CPU PAR	min	1.631	max	1.776	mean	1.709	839	mhd4800a.mtx							
766		H	min	11.873	max	11.877	mean	11.875	840	Regular							
767	OPF_3754.mtx								841		CPU COO	min	0.759	max	0.795	mean	0.780
768	Regular								842		CPU CSR	min	2.479	max	2.565	mean	2.557
769		CPU COO	min	0.726	max	0.774	mean	0.747	843		GPU 64 COO	min	5.490	max	5.650	mean	5.552
770		CPU CSR	min	2.898	max	2.919	mean	2.908	844		CSR	min	16.700	max	19.460	mean	18.004
771		GPU 64 COO	min	7.600	max	7.820	mean	7.766	845		CPU PAR	min	1.456	max	1.523	mean	1.492
772		CSR	min	25.070	max	29.030	mean	26.756	846		H	min	7.132	max	7.132	mean	7.132
773		CPU PAR	min	1.437	max	1.508	mean	1.471	847	Row-Premute							
774		H	min	8.393	max	8.393	mean	8.393	848		CPU COO	min	0.695	max	0.943	mean	0.726
775	Row-Premute								849		CPU CSR	min	2.480	max	2.488	mean	2.485
776		CPU COO	min	0.714	max	1.574	mean	0.817	850		GPU 64 COO	min	5.410	max	5.490	mean	5.453
777		CPU CSR	min	2.686	max	2.711	mean	2.699	851		CSR	min	15.700	max	17.520	mean	16.678
778		GPU 64 COO	min	7.410	max	7.570	mean	7.484	852		CPU PAR	min	1.422	max	1.514	mean	1.474
779		CSR	min	19.600	max	21.190	mean	20.307	853		H	min	10.959	max	10.966	mean	10.963
780		CPU PAR	min	1.443	max	1.505	mean	1.469	854	Row-Gradient							
781		H	min	11.267	max	11.272	mean	11.269	855		CPU COO	min	0.723	max	2.029	mean	0.990
782	Row-Gradient								856		CPU CSR	min	2.411	max	2.427	mean	2.421
783		CPU COO	min	0.723	max	1.232	mean	0.775	857		GPU 64 COO	min	5.490	max	5.560	mean	5.534
784		CPU CSR	min	1.672	max	1.691	mean	1.685	858		CSR	min	16.350	max	19.560	mean	17.784
785		GPU 64 COO	min	7.600	max	7.760	mean	7.716	859		CPU PAR	min	1.441	max	1.509	mean	1.477
786		CSR	min	23.160	max	25.590	mean	24.304	860		H	min	9.512	max	9.526	mean	9.520
787		CPU PAR	min	1.675	max	1.736	mean	1.703	861	Column-Gradient							
788		H	min	10.463	max	10.472	mean	10.468	862		CPU COO	min	0.721	max	1.802	mean	0.871
789	Column-Gradient								863		CPU CSR	min	2.393	max	2.408	mean	2.404
790		CPU COO	min	0.726	max	1.431	mean	0.778	864		GPU 64 COO	min	5.410	max	5.480	mean	5.453
791		CPU CSR	min	1.671	max	1.685	mean	1.679	865		CSR	min	15.680	max	17.870	mean	16.540
792		GPU 64 COO	min	7.410	max	7.530	mean	7.467	866		CPU PAR	min	1.429	max	1.488	mean	1.468
793		CSR	min	18.140	max	20.350	mean	19.315	867		H	min	10.931	max	10.945	mean	10.938
794		CPU PAR	min	1.650	max	1.736	mean	1.699	868	Row-Column-Permute							
795		H	min	11.393	max	11.401	mean	11.397	869		CPU COO	min	0.728	max	1.646	mean	1.037
796	Row-Column-Permute								870		CPU CSR	min	2.472	max	2.488	mean	2.480
797		CPU COO	min	0.711	max	1.458	mean	0.751	871		GPU 64 COO	min	5.410	max	5.480	mean	5.449
798		CPU CSR	min	2.678	max	2.717	mean	2.700	872		CSR	min	15.760	max	17.560	mean	16.654
799		GPU 64 COO	min	7.400	max	7.540	mean	7.471	873		CPU PAR	min	1.428	max	1.513	mean	1.474
800		CSR	min	19.560	max	21.150	mean	20.453	874		H	min	10.959	max	10.967	mean	10.963
801		CPU PAR	min	1.440	max	1.499	mean	1.467	875	gen4.mtx							
802		H	min	11.266	max	11.272	mean	11.269	876	Regular							
803	c-47.mtx								877		CPU COO	min	0.737	max	1.977	mean	1.431
804	Regular								878		CPU CSR	min	2.674	max	2.688	mean	2.681
805		CPU COO	min	0.754	max	1.829	mean	1.204	879		GPU 64 COO	min	5.900	max	6.000	mean	5.954
806		CPU CSR	min	2.610	max	2.624	mean	2.618	880		CSR	min	13.650	max	15.410	mean	14.657
807		GPU 64 COO	min	9.530	max	9.870	mean	9.640	881		CPU PAR	min	1.468	max	1.521	mean	1.491
808		CSR	min	23.990	max	25.910	mean	24.992	882		H	min	9.234	max	9.234	mean	9.234
809		CPU PAR	min	1.311	max	1.380	mean	1.357	883	Row-Premute							
810		H	min	8.364	max	8.364	mean	8.364	884		CPU COO	min	0.740	max	2.048	mean	1.121
811	Row-Premute								885		CPU CSR	min	2.777	max	2.798	mean	2.790
812		CPU COO	min	0.740	max	0.885	mean	0.755	886		GPU 64 COO	min	5.910	max	5.970	mean	5.944
813		CPU CSR	min	2.574	max	2.611	mean	2.597	887		CSR	min	13.700	max	15.370	mean	14.541
814		GPU 64 COO	min	9.320	max	9.510	mean	9.397	888		CPU PAR	min	1.468	max	1.546	mean	1.502

889		H	min 10.250 max 10.255 mean 10.252	963		CPU COO	min 0.735 max 1.806 mean 0.878
890	Row-Gradient			964		CPU CSR	min 2.706 max 2.744 mean 2.726
891		CPU COO	min 0.740 max 1.790 mean 0.994	965		GPU 64 COO	min 6.390 max 6.500 mean 6.433
892		CPU CSR	min 2.663 max 2.682 mean 2.674	966		CSR	min 19.780 max 22.870 mean 20.936
893		GPU 64 COO	min 5.890 max* 6.160 mean 5.946	967		CPU PAR	min 1.710 max 1.865 mean 1.785
894		CSR	min 13.780 max*17.520 mean 15.601	968		H	min 10.251 max 10.267 mean 10.257
895		CPU PAR	min 1.479 max* 1.619 mean 1.569	969	Column-Gradient		
896		H	min 9.939 max 9.955 mean 9.948	970		CPU COO	min 0.728 max 1.792 mean 0.986
897	Column-Gradient			971		CPU CSR	min 2.521 max 2.720 mean 2.703
898		CPU COO	min 0.743 max 1.991 mean 0.981	972		GPU 64 COO	min 6.280 max 6.370 mean 6.327
899		CPU CSR	min 2.620 max 2.654 mean 2.646	973		CSR	min 18.000 max 19.720 mean 19.040
900		GPU 64 COO	min 5.840 max 5.910 mean 5.885	974		CPU PAR	min 1.649 max 1.741 mean 1.702
901		CSR	min 13.130 max 17.040 mean 15.008	975		H	min 11.113 max 11.121 mean 11.117
902		CPU PAR	min 1.477 max 1.607 mean 1.559	976	Row-Column-Permute		
903		H	min 10.858 max*10.876 mean 10.864	977		CPU COO	min 0.714 max 1.525 mean 0.957
904	Row-Column-Permute			978		CPU CSR	min 2.876 max 2.892 mean 2.884
905		CPU COO	min 0.742 max 2.010 mean 1.124	979		GPU 64 COO	min 6.280 max 6.370 mean 6.322
906		CPU CSR	min 2.789 max* 2.800 mean 2.795	980		CSR	min 17.960 max 19.670 mean 18.670
907		GPU 64 COO	min 5.900 max 5.980 mean 5.941	981		CPU PAR	min 1.667 max 1.754 mean 1.710
908		CSR	min 13.640 max 15.410 mean 14.556	982		H	min 11.162 max*11.168 mean 11.165
909		CPU PAR	min 1.462 max 1.540 mean 1.504	983	TSOPF_RS_b39_c7.mtx		
910		H	min 10.250 max 10.253 mean 10.252	984	Regular		
911	Maragal_6.mtx			985		CPU COO	min 0.771 max 0.793 mean 0.780
912	Regular			986		CPU CSR	min 3.219 max* 3.232 mean 3.227
913		CPU COO	min 0.725 max 0.741 mean 0.729	987		GPU 64 COO	min 11.070 max*11.200 mean 11.142
914		CPU CSR	min 2.345 max 2.409 mean 2.372	988		CSR	min 37.050 max*42.100 mean 39.040
915		GPU 64 COO	min 18.200 max 18.770 mean 18.357	989		CPU PAR	min 1.910 max 2.027 mean 1.982
916		CSR	min 38.310 max*40.240 mean 39.477	990		H	min 7.304 max 7.304 mean 7.304
917		CPU PAR	min 0.789 max 0.813 mean 0.797	991	Row-Premute		
918		H	min 9.930 max 9.930 mean 9.930	992		CPU COO	min 0.701 max 0.722 mean 0.707
919	Row-Premute			993		CPU CSR	min 2.931 max 2.952 mean 2.942
920		CPU COO	min 0.709 max 0.779 mean 0.715	994		GPU 64 COO	min 10.860 max 11.030 mean 10.928
921		CPU CSR	min 2.675 max 2.715 mean 2.696	995		CSR	min 28.730 max 30.880 mean 29.483
922		GPU 64 COO	min 17.810 max 18.030 mean 17.935	996		CPU PAR	min 1.760 max 1.922 mean 1.851
923		CSR	min 29.650 max 30.580 mean 30.109	997		H	min 10.537 max 10.541 mean 10.539
924		CPU PAR	min 0.857 max 0.940 mean 0.904	998	Row-Gradient		
925		H	min 10.777 max 10.779 mean 10.778	999		CPU COO	min 0.747 max 0.808 mean 0.757
926	Row-Gradient			1000		CPU CSR	min 2.606 max 2.648 mean 2.624
927		CPU COO	min 0.710 max* 1.566 mean 0.755	1001		GPU 64 COO	min 10.850 max 11.120 mean 10.999
928		CPU CSR	min 2.042 max 2.159 mean 2.120	1002		CSR	min 33.910 max 37.600 mean 35.909
929		GPU 64 COO	min 18.460 max*18.960 mean 18.665	1003		CPU PAR	min 2.154 max* 2.245 mean 2.203
930		CSR	min 25.650 max 27.330 mean 26.549	1004		H	min 9.636 max 9.646 mean 9.642
931		CPU PAR	min 2.257 max 2.612 mean 2.416	1005	Column-Gradient		
932		H	min 11.251 max 11.301 mean 11.285	1006		CPU COO	min 0.718 max* 1.693 mean 0.802
933	Column-Gradient			1007		CPU CSR	min 2.502 max 2.585 mean 2.547
934		CPU COO	min 0.711 max 0.743 mean 0.725	1008		GPU 64 COO	min 10.700 max 10.990 mean 10.804
935		CPU CSR	min 2.036 max 2.161 mean 2.110	1009		CSR	min 27.230 max 29.380 mean 28.488
936		GPU 64 COO	min 17.840 max 18.860 mean 18.149	1010		CPU PAR	min 2.128 max 2.227 mean 2.172
937		CSR	min 19.410 max 20.690 mean 20.066	1011		H	min 11.131 max*11.222 mean 11.208
938		CPU PAR	min 2.174 max* 2.546 mean 2.349	1012	Row-Column-Permute		
939		H	min 12.011 max*12.072 mean 12.052	1013		CPU COO	min 0.709 max 0.726 mean 0.716
940	Row-Column-Permute			1014		CPU CSR	min 2.917 max 2.958 mean 2.940
941		CPU COO	min 0.712 max 0.971 mean 0.737	1015		GPU 64 COO	min 10.840 max 11.030 mean 10.930
942		CPU CSR	min 2.732 max* 2.751 mean 2.743	1016		CSR	min 28.780 max 30.810 mean 29.578
943		GPU 64 COO	min 17.720 max 18.070 mean 17.911	1017		CPU PAR	min 1.757 max 1.834 mean 1.792
944		CSR	min 29.600 max 30.500 mean 29.961	1018		H	min 10.537 max 10.540 mean 10.539
945		CPU PAR	min 0.827 max 0.954 mean 0.913				
946		H	min 10.776 max 10.778 mean 10.777				
947	aft01.mtx			1019	10 ELLESMERE		
948	Regular			1020	aft01.mtx		
949		CPU COO	min 0.735 max* 2.079 mean 1.069	1021	Regular		
950		CPU CSR	min 3.132 max* 3.154 mean 3.145	1022		GPU 64 COO	min 4.080 max* 4.280 mean 4.186
951		GPU 64 COO	min 6.390 max* 6.610 mean 6.457	1023		CSR	min 9.660 max*12.660 mean 11.485
952		CSR	min 19.990 max*23.250 mean 21.820	1024		H	min 7.811 max 7.811 mean 7.811
953		CPU PAR	min 1.746 max* 1.865 mean 1.812	1025	Row-Premute		
954		H	min 7.811 max 7.811 mean 7.811	1026		GPU 64 COO	min 3.860 max 4.090 mean 4.001
955	Row-Premute			1027		CSR	min 9.520 max 10.340 mean 9.936
956		CPU COO	min 0.714 max 1.648 mean 0.840	1028		H	min 11.161 max 11.167 mean 11.165
957		CPU CSR	min 2.864 max 2.892 mean 2.883	1029	Row-Gradient		
958		GPU 64 COO	min 6.280 max 6.380 mean 6.329	1030		GPU 64 COO	min 4.010 max 4.240 mean 4.135
959		CSR	min 17.980 max 19.700 mean 19.105	1031		CSR	min 5.890 max 11.350 mean 6.882
960		CPU PAR	min 1.729 max 1.850 mean 1.782	1032		H	min 10.246 max 10.262 mean 10.256
961		H	min 11.162 max 11.168 mean 11.165	1033	Column-Gradient		
962	Row-Gradient						

1034		GPU 64 COO min 3.850 max 4.100 mean 4.012	1108		H min 7.380 max 7.380 mean 7.380
1035		CSR min 5.460 max 8.790 mean 6.005	1109	Row-Premute	
1036		H min 11.112 max 11.122 mean 11.117	1110		GPU 64 COO min 4.820 max 4.940 mean 4.859
1037	Row-Column-Permute		1111		CSR min 5.080 max 6.520 mean 6.342
1038		GPU 64 COO min 3.850 max 4.080 mean 3.990	1112		H min 10.042 max 10.047 mean 10.044
1039		CSR min 5.420 max 6.760 mean 5.977	1113	Row-Gradient	
1040		H min 11.162 max*11.169 mean 11.165	1114		GPU 64 COO min 4.810 max* 4.940 mean 4.876
1041	bloweya.mtx		1115		CSR min 6.100 max* 6.560 mean 6.307
1042	Regular		1116		H min 9.681 max 9.704 mean 9.694
1043		GPU 64 COO min 0.000 max 0.000 mean 0.000	1117	Column-Gradient	
1044		CSR min 0.000 max 0.000 mean 0.000	1118		GPU 64 COO min 4.810 max 4.930 mean 4.869
1045		H min 7.205 max 7.205 mean 7.205	1119		CSR min 4.820 max 6.460 mean 6.208
1046	Row-Premute		1120		H min 10.554 max*10.661 mean 10.638
1047		GPU 64 COO min 3.800 max 3.940 mean 3.875	1121	Row-Column-Permute	
1048		CSR min 3.710 max 4.570 mean 4.399	1122		GPU 64 COO min 4.810 max 4.940 mean 4.864
1049		H min 11.025 max 11.031 mean 11.028	1123		CSR min 5.930 max 6.520 mean 6.379
1050	Row-Gradient		1124		H min 10.041 max 10.047 mean 10.044
1051		GPU 64 COO min 3.800 max* 4.120 mean 3.962	1125	cvxqp3.mtx	
1052		CSR min 4.340 max* 4.670 mean 4.546	1126	Regular	
1053		H min 10.296 max 10.307 mean 10.300	1127		GPU 64 COO min 3.350 max* 3.590 mean 3.483
1054	Column-Gradient		1128		CSR min 5.430 max* 9.260 mean 8.333
1055		GPU 64 COO min 3.880 max 4.100 mean 3.978	1129		H min 8.646 max 8.646 mean 8.646
1056		CSR min 4.240 max 4.570 mean 4.412	1130	Row-Premute	
1057		H min 10.881 max 10.886 mean 10.883	1131		GPU 64 COO min 3.230 max 3.480 mean 3.371
1058	Row-Column-Permute		1132		CSR min 7.560 max 8.220 mean 7.900
1059		GPU 64 COO min 3.800 max 3.980 mean 3.885	1133		H min 11.027 max 11.033 mean 11.030
1060		CSR min 4.130 max 4.540 mean 4.399	1134	Row-Gradient	
1061		H min 11.025 max*11.033 mean 11.029	1135		GPU 64 COO min 3.240 max 3.510 mean 3.396
1062	brainpc2.mtx		1136		CSR min 6.990 max 7.890 mean 7.574
1063	Regular		1137		H min 11.060 max 11.069 mean 11.064
1064		GPU 64 COO min 0.000 max 0.000 mean 0.000	1138	Column-Gradient	
1065		CSR min 0.000 max 0.000 mean 0.000	1139		GPU 64 COO min 3.240 max 3.480 mean 3.374
1066		H min 7.478 max 7.478 mean 7.478	1140		CSR min 6.980 max 7.900 mean 7.557
1067	Row-Premute		1141		H min 11.126 max*11.134 mean 11.130
1068		GPU 64 COO min 3.840 max* 6.750 mean 4.110	1142	Row-Column-Permute	
1069		CSR min 4.260 max* 4.500 mean 4.437	1143		GPU 64 COO min 3.110 max 3.470 mean 3.365
1070		H min 9.809 max 9.813 mean 9.811	1144		CSR min 4.810 max 8.210 mean 7.742
1071	Row-Gradient		1145		H min 11.026 max 11.032 mean 11.030
1072		GPU 64 COO min 0.640 max 4.030 mean 3.864	1146	ex19.mtx	
1073		CSR min 4.270 max 4.470 mean 4.383	1147	Regular	
1074		H min 9.722 max 9.727 mean 9.724	1148		GPU 64 COO min 2.450 max* 2.610 mean 2.564
1075	Column-Gradient		1149		CSR min 4.490 max 4.760 mean 4.714
1076		GPU 64 COO min 0.640 max 4.070 mean 3.898	1150		H min 8.228 max 8.228 mean 8.228
1077		CSR min 4.230 max 4.500 mean 4.386	1151	Row-Premute	
1078		H min 10.368 max*10.372 mean 10.370	1152		GPU 64 COO min 2.000 max 2.040 mean 2.021
1079	Row-Column-Permute		1153		CSR min 4.640 max 4.780 mean 4.733
1080		GPU 64 COO min 3.980 max 4.110 mean 4.027	1154		H min 11.835 max 11.840 mean 11.838
1081		CSR min 4.320 max 4.490 mean 4.437	1155	Row-Gradient	
1082		H min 9.809 max 9.813 mean 9.811	1156		GPU 64 COO min 2.240 max 2.390 mean 2.329
1083	c-47.mtx		1157		CSR min 4.570 max* 4.850 mean 4.807
1084	Regular		1158		H min 10.742 max 10.752 mean 10.747
1085		GPU 64 COO min 3.980 max* 4.080 mean 4.026	1159	Column-Gradient	
1086		CSR min 4.760 max 4.850 mean 4.812	1160		GPU 64 COO min 2.010 max 2.050 mean 2.034
1087		H min 8.364 max 8.364 mean 8.364	1161		CSR min 4.570 max 4.760 mean 4.701
1088	Row-Premute		1162		H min 11.872 max*11.881 mean 11.878
1089		GPU 64 COO min 3.880 max 4.010 mean 3.942	1163	Row-Column-Permute	
1090		CSR min 4.040 max 4.900 mean 4.807	1164		GPU 64 COO min 2.000 max 2.040 mean 2.023
1091		H min 10.059 max 10.063 mean 10.061	1165		CSR min 0.770 max 4.780 mean 4.594
1092	Row-Gradient		1166		H min 11.835 max 11.840 mean 11.838
1093		GPU 64 COO min 3.900 max 4.050 mean 3.976	1167	gen4.mtx	
1094		CSR min 4.380 max 4.740 mean 4.630	1168	Regular	
1095		H min 10.201 max 10.228 mean 10.214	1169		GPU 64 COO min 4.880 max 4.980 mean 4.900
1096	Column-Gradient		1170		CSR min 10.020 max*11.300 mean 10.716
1097		GPU 64 COO min 3.860 max 3.990 mean 3.936	1171		H min 9.234 max 9.234 mean 9.234
1098		CSR min 4.350 max 4.610 mean 4.525	1172	Row-Premute	
1099		H min 11.204 max*11.241 mean 11.222	1173		GPU 64 COO min 4.860 max 4.930 mean 4.890
1100	Row-Column-Permute		1174		CSR min 0.330 max 11.200 mean 10.038
1101		GPU 64 COO min 3.890 max 4.020 mean 3.953	1175		H min 10.249 max 10.254 mean 10.252
1102		CSR min 4.490 max* 4.920 mean 4.840	1176	Row-Gradient	
1103		H min 10.058 max 10.063 mean 10.061	1177		GPU 64 COO min 4.860 max* 4.990 mean 4.908
1104	case9.mtx		1178		CSR min 9.160 max 11.240 mean 10.435
1105	Regular		1179		H min 9.939 max 9.961 mean 9.947
1106		GPU 64 COO min 0.000 max 0.000 mean 0.000	1180	Column-Gradient	
1107		CSR min 0.000 max 0.000 mean 0.000	1181		GPU 64 COO min 4.780 max 4.880 mean 4.816

1182		CSR min 7.770 max 10.570 mean 9.407	1256	Row-Permute	
1183		H min 10.851 max*10.876 mean 10.864	1257		GPU 64 COO min 4.420 max 4.520 mean 4.445
1184	Row-Column-Permute		1258		CSR min 10.520 max 10.880 mean 10.696
1185		GPU 64 COO min 4.850 max 4.950 mean 4.886	1259		H min 10.960 max*10.968 mean 10.963
1186		CSR min 10.220 max 11.280 mean 10.748	1260	Row-Gradient	
1187		H min 10.250 max 10.255 mean 10.252	1261		GPU 64 COO min 4.570 max 4.690 mean 4.605
1188	lp_fit2d.mtx		1262		CSR min 4.550 max 13.350 mean 12.479
1189	Regular		1263		H min 9.508 max 9.527 mean 9.520
1190		GPU 64 COO min 4.360 max* 4.640 mean 4.515	1264	Column-Gradient	
1191		CSR min 10.080 max 10.900 mean 10.491	1265		GPU 64 COO min 4.430 max 4.530 mean 4.461
1192		H min 11.109 max 11.109 mean 11.109	1266		CSR min 10.250 max 10.940 mean 10.603
1193	Row-Permute		1267		H min 10.934 max 10.945 mean 10.939
1194		GPU 64 COO min 4.170 max 4.630 mean 4.476	1268	Row-Column-Permute	
1195		CSR min 0.910 max 10.910 mean 10.257	1269		GPU 64 COO min 4.420 max 4.520 mean 4.450
1196		H min 11.098 max 11.104 mean 11.101	1270		CSR min 7.380 max 10.900 mean 10.598
1197	Row-Gradient		1271		H min 10.959 max 10.967 mean 10.963
1198		GPU 64 COO min 4.370 max 4.630 mean 4.529	1272	mult_dcop_01.mtx	
1199		CSR min 10.030 max 10.970 mean 10.624	1273	Regular	
1200		H min 11.109 max 11.109 mean 11.109	1274		GPU 64 COO min 3.420 max 3.630 mean 3.555
1201	Column-Gradient		1275		CSR min 3.650 max 4.090 mean 3.814
1202		GPU 64 COO min 4.250 max 4.640 mean 4.499	1276		H min 9.689 max 9.689 mean 9.689
1203		CSR min 8.510 max*11.010 mean 10.505	1277	Row-Permute	
1204		H min 11.328 max*11.333 mean 11.331	1278		GPU 64 COO min 3.450 max 3.580 mean 3.521
1205	Row-Column-Permute		1279		CSR min 3.610 max 4.150 mean 3.785
1206		GPU 64 COO min 4.350 max 4.640 mean 4.511	1280		H min 10.738 max 10.742 mean 10.740
1207		CSR min 10.040 max 10.790 mean 10.468	1281	Row-Gradient	
1208		H min 11.097 max 11.106 mean 11.101	1282		GPU 64 COO min 3.510 max* 3.660 mean 3.579
1209	lp_osa_07.mtx		1283		CSR min 3.650 max 4.160 mean 3.806
1210	Regular		1284		H min 10.576 max 10.585 mean 10.580
1211		GPU 64 COO min 0.460 max* 3.640 mean 3.456	1285	Column-Gradient	
1212		CSR min 5.570 max* 8.530 mean 8.106	1286		GPU 64 COO min 3.460 max 3.650 mean 3.584
1213		H min 8.412 max 8.412 mean 8.412	1287		CSR min 3.660 max* 4.240 mean 3.799
1214	Row-Permute		1288		H min 10.826 max*10.842 mean 10.836
1215		GPU 64 COO min 3.140 max 3.450 mean 3.367	1289	Row-Column-Permute	
1216		CSR min 7.600 max 8.070 mean 7.853	1290		GPU 64 COO min 3.470 max 3.580 mean 3.532
1217		H min 9.255 max 9.258 mean 9.256	1291		CSR min 3.600 max 3.980 mean 3.743
1218	Row-Gradient		1292		H min 10.738 max 10.742 mean 10.740
1219		GPU 64 COO min 3.190 max 3.610 mean 3.509	1293	mult_dcop_02.mtx	
1220		CSR min 0.000 max 8.260 mean 7.597	1294	Regular	
1221		H min 8.583 max 8.678 mean 8.670	1295		GPU 64 COO min 3.390 max 3.660 mean 3.585
1222	Column-Gradient		1296		CSR min 0.960 max 4.330 mean 4.162
1223		GPU 64 COO min 3.330 max 3.500 mean 3.416	1297		H min 9.689 max 9.689 mean 9.689
1224		CSR min 6.730 max 7.540 mean 7.199	1298	Row-Permute	
1225		H min 9.542 max* 9.604 mean 9.581	1299		GPU 64 COO min 3.310 max 3.600 mean 3.488
1226	Row-Column-Permute		1300		CSR min 0.620 max 4.290 mean 4.132
1227		GPU 64 COO min 3.290 max 3.430 mean 3.365	1301		H min 10.738 max 10.743 mean 10.740
1228		CSR min 7.390 max 8.060 mean 7.832	1302	Row-Gradient	
1229		H min 9.255 max 9.258 mean 9.256	1303		GPU 64 COO min 3.310 max* 3.670 mean 3.593
1230	Maragal_6.mtx		1304		CSR min 4.130 max* 4.430 mean 4.331
1231	Regular		1305		H min 10.576 max 10.584 mean 10.580
1232		GPU 64 COO min 4.160 max 4.310 mean 4.217	1306	Column-Gradient	
1233		CSR min 4.940 max 4.960 mean 4.956	1307		GPU 64 COO min 0.550 max 3.660 mean 3.486
1234		H min 9.930 max 9.930 mean 9.930	1308		CSR min 3.890 max 4.410 mean 4.275
1235	Row-Permute		1309		H min 10.831 max*10.843 mean 10.836
1236		GPU 64 COO min 4.220 max 4.240 mean 4.225	1310	Row-Column-Permute	
1237		CSR min 4.750 max*13.040 mean 5.133	1311		GPU 64 COO min 3.470 max 3.590 mean 3.542
1238		H min 10.776 max 10.778 mean 10.777	1312		CSR min 4.190 max 4.290 mean 4.242
1239	Row-Gradient		1313		H min 10.738 max 10.742 mean 10.740
1240		GPU 64 COO min 4.180 max* 4.450 mean 4.245	1314	mult_dcop_03.mtx	
1241		CSR min 4.880 max 4.940 mean 4.915	1315	Regular	
1242		H min 11.259 max*11.302 mean 11.281	1316		GPU 64 COO min 3.360 max* 3.660 mean 3.550
1243	Column-Gradient		1317		CSR min 3.650 max 4.090 mean 3.813
1244		GPU 64 COO min 4.200 max 4.250 mean 4.236	1318		H min 9.689 max 9.689 mean 9.689
1245		CSR min 4.800 max 4.890 mean 4.859	1319	Row-Permute	
1246		H min 12.022 max 12.073 mean 12.051	1320		GPU 64 COO min 3.450 max 3.580 mean 3.521
1247	Row-Column-Permute		1321		CSR min 3.610 max 4.160 mean 3.784
1248		GPU 64 COO min 4.210 max 4.230 mean 4.222	1322		H min 10.738 max 10.743 mean 10.740
1249		CSR min 4.860 max 4.890 mean 4.887	1323	Row-Gradient	
1250		H min 10.776 max 10.778 mean 10.778	1324		GPU 64 COO min 3.470 max 3.660 mean 3.572
1251	mhd4800a.mtx		1325		CSR min 3.640 max 4.190 mean 3.809
1252	Regular		1326		H min 10.572 max 10.584 mean 10.580
1253		GPU 64 COO min 4.570 max* 4.710 mean 4.608	1327	Column-Gradient	
1254		CSR min 12.690 max*13.940 mean 13.369	1328		GPU 64 COO min 3.430 max 3.650 mean 3.562
1255		H min 7.132 max 7.132 mean 7.132	1329		CSR min 3.670 max* 4.290 mean 3.793

1330		H	min 10.828 max*10.840 mean 10.834	1404		GPU 64 COO min 4.540 max 4.940 mean 4.874
1331	Row-Column-Permute			1405		CSR min 6.280 max 6.520 mean 6.403
1332		GPU 64 COO min 3.370 max 3.610 mean 3.502		1406	H	min 10.042 max 10.047 mean 10.044
1333		CSR min 3.610 max 3.970 mean 3.744		1407	Row-Gradient	
1334		H min 10.738 max 10.741 mean 10.740		1408		GPU 64 COO min 4.830 max 4.930 mean 4.875
1335	OPF_3754.mtx			1409		CSR min 5.790 max* 6.560 mean 6.289
1336	Regular			1410	H	min 9.675 max 9.706 mean 9.692
1337		GPU 64 COO min 4.700 max* 4.930 mean 4.842		1411	Column-Gradient	
1338		CSR min 6.230 max* 6.600 mean 6.411		1412		GPU 64 COO min 4.790 max* 4.960 mean 4.880
1339		H min 8.393 max 8.393 mean 8.393		1413		CSR min 5.760 max 6.450 mean 6.204
1340	Row-Premute			1414	H	min 10.601 max*10.661 mean 10.626
1341		GPU 64 COO min 4.620 max 4.890 mean 4.787		1415	Row-Column-Permute	
1342		CSR min 5.780 max 6.310 mean 6.192		1416		GPU 64 COO min 4.330 max 4.950 mean 4.845
1343		H min 11.265 max 11.272 mean 11.269		1417		CSR min 5.740 max 6.500 mean 6.375
1344	Row-Gradient			1418		H min 10.041 max 10.046 mean 10.044
1345		GPU 64 COO min 4.570 max 4.870 mean 4.776		1419	TSOPF_RS_b39_c7.mtx	
1346		CSR min 5.770 max 6.510 mean 6.302		1420	Regular	
1347		H min 10.464 max 10.473 mean 10.468		1421		GPU 64 COO min 4.300 max* 4.430 mean 4.364
1348	Column-Gradient			1422		CSR min 4.480 max 4.750 mean 4.716
1349		GPU 64 COO min 4.580 max 4.870 mean 4.756		1423	H	min 7.304 max 7.304 mean 7.304
1350		CSR min 5.630 max 6.180 mean 6.055		1424	Row-Premute	
1351		H min 11.394 max*11.401 mean 11.397		1425		GPU 64 COO min 4.260 max 4.400 mean 4.353
1352	Row-Column-Permute			1426		CSR min 4.490 max 4.770 mean 4.734
1353		GPU 64 COO min 4.610 max 4.900 mean 4.780		1427	H	min 10.536 max 10.541 mean 10.539
1354		CSR min 5.010 max 6.300 mean 6.113		1428	Row-Gradient	
1355		H min 11.268 max 11.272 mean 11.270		1429		GPU 64 COO min 3.970 max 4.420 mean 4.338
1356	OPF_6000.mtx			1430		CSR min 4.620 max* 4.820 mean 4.789
1357	Regular			1431	H	min 9.638 max 9.644 mean 9.641
1358		GPU 64 COO min 3.780 max* 3.920 mean 3.864		1432	Column-Gradient	
1359		CSR min 4.270 max 4.360 mean 4.332		1433		GPU 64 COO min 4.240 max 4.430 mean 4.368
1360		H min 8.799 max 8.799 mean 8.799		1434		CSR min 4.710 max 4.770 mean 4.736
1361	Row-Premute			1435	H	min 11.129 max*11.222 mean 11.205
1362		GPU 64 COO min 3.770 max 3.870 mean 3.821		1436	Row-Column-Permute	
1363		CSR min 3.970 max*11.050 mean 4.439		1437		GPU 64 COO min 4.260 max 4.410 mean 4.359
1364		H min 11.872 max 11.877 mean 11.875		1438		CSR min 4.660 max 4.760 mean 4.738
1365	Row-Gradient			1439	H	min 10.537 max 10.541 mean 10.539
1366		GPU 64 COO min 3.700 max 3.870 mean 3.795				
1367		CSR min 4.330 max 4.440 mean 4.403				
1368		H min 11.109 max 11.116 mean 11.113				
1369	Column-Gradient			1440	11 FIJI	
1370		GPU 64 COO min 3.690 max 3.870 mean 3.804		1441	mult_dcop_03.mtx	
1371		CSR min 4.260 max 4.340 mean 4.308		1442	Regular	
1372		H min 12.041 max*12.045 mean 12.043		1443		GPU 64 COO min 5.140 max* 5.140 mean 5.140
1373	Row-Column-Permute			1444		CSR min 10.340 max*10.390 mean 10.365
1374		GPU 64 COO min 3.780 max 3.860 mean 3.819		1445	H	min 9.689 max 9.689 mean 9.689
1375		CSR min 4.090 max 4.290 mean 4.259		1446	Row-Premute	
1376		H min 11.873 max 11.877 mean 11.876		1447		GPU 64 COO min 4.970 max 4.990 mean 4.980
1377	shermanACb.mtx			1448		CSR min 9.420 max 9.430 mean 9.425
1378	Regular			1449	H	min 10.739 max 10.739 mean 10.739
1379		GPU 64 COO min 2.920 max* 3.140 mean 3.048		1450	Row-Gradient	
1380		CSR min 5.550 max 5.980 mean 5.803		1451		GPU 64 COO min 5.080 max 5.090 mean 5.085
1381		H min 8.600 max 8.600 mean 8.600		1452		CSR min 9.720 max 10.300 mean 10.010
1382	Row-Premute			1453	H	min 10.579 max 10.582 mean 10.580
1383		GPU 64 COO min 2.760 max 3.020 mean 2.898		1454	Column-Gradient	
1384		CSR min 2.660 max 5.830 mean 5.632		1455		GPU 64 COO min 5.030 max 5.120 mean 5.075
1385		H min 10.377 max 10.381 mean 10.379		1456		CSR min 9.330 max 9.770 mean 9.550
1386	Row-Gradient			1457	H	min 10.835 max*10.838 mean 10.836
1387		GPU 64 COO min 2.800 max 3.040 mean 2.944		1458	Row-Column-Permute	
1388		CSR min 5.330 max* 6.020 mean 5.742		1459		GPU 64 COO min 5.000 max 5.010 mean 5.005
1389		H min 9.919 max 9.925 mean 9.922		1460		CSR min 7.580 max 9.460 mean 8.520
1390	Column-Gradient			1461	H	min 10.739 max 10.741 mean 10.740
1391		GPU 64 COO min 2.720 max 3.010 mean 2.926		1462	mult_dcop_03.mtx	
1392		CSR min 0.000 max 5.840 mean 5.513		1463	Regular	
1393		H min 10.587 max*10.596 mean 10.591		1464		GPU 64 COO min 5.140 max* 5.140 mean 5.140
1394	Row-Column-Permute			1465		CSR min 10.340 max*10.390 mean 10.365
1395		GPU 64 COO min 2.780 max 3.030 mean 2.939		1466	H	min 9.689 max 9.689 mean 9.689
1396		CSR min 4.860 max 5.810 mean 5.667		1467	Row-Premute	
1397		H min 10.376 max 10.382 mean 10.379		1468		GPU 64 COO min 4.970 max 4.990 mean 4.980
1398	TSOPF_FS_b9_c6.mtx			1469		CSR min 9.420 max 9.430 mean 9.425
1399	Regular			1470	H	min 10.739 max 10.739 mean 10.739
1400		GPU 64 COO min 0.000 max 0.000 mean 0.000		1471	Row-Gradient	
1401		CSR min 0.000 max 0.000 mean 0.000		1472		GPU 64 COO min 5.080 max 5.090 mean 5.085
1402		H min 7.380 max 7.380 mean 7.380		1473		CSR min 9.720 max 10.300 mean 10.010
1403	Row-Premute			1474	H	min 10.579 max 10.582 mean 10.580

1475	Column-Gradient					1549				CSR min	6.360	max	7.450	mean	6.711					
1476		GPU 64	COO	min	5.030	max	5.120	mean	5.075											
1477				CSR min	9.330	max	9.770	mean	9.550	1551	Row-Premute									
1478		H		min	10.835	max	10.838	mean	10.836	1552		GPU 64	COO	min	3.950					
1479	Row-Column-Permute					1553				1554			max	3.980	mean	3.953				
1480		GPU 64	COO	min	5.000	max	5.010	mean	5.005	1555			CSR min	6.330	max	7.400				
1481				CSR min	7.580	max	9.460	mean	8.520	1556		H		min	11.098	max	11.104			
1482		H		min	10.739	max	10.741	mean	10.740	1557	Row-Gradient				mean	11.101				
1483	mult_dcop_03.mtx					1558				1559		GPU 64	COO	min	3.960	max	3.980			
1484	Regular					1559	Column-Gradient			1560			CSR min	6.270	max	10.770	mean	7.017		
1485		GPU 64	COO	min	5.130	max	5.220	mean	5.142	1561		H		min	11.109	max	11.109	mean	11.109	
1486				CSR min	7.250	max	9.320	mean	7.722	1562		GPU 64	COO	min	3.940	max	3.960	mean	3.950	
1487		H		min	9.689	max	9.689	mean	9.689	1563	Row-Column-Permute			CSR min	6.270	max	7.370	mean	6.696	
1488	Row-Premute					1564				1565		H		min	11.329	max	11.334	mean	11.333	
1489		GPU 64	COO	min	4.980	max	5.030	mean	4.999	1566		GPU 64	COO	min	3.950	max	3.960	mean	3.952	
1490				CSR min	6.460	max	8.470	mean	6.950	1567			CSR min	6.180	max	7.420	mean	6.641		
1491		H		min	10.738	max	10.742	mean	10.740	1568	Row-Gradient		H		min	11.098	max	11.105	mean	11.101
1492	Row-Gradient					1569				1570			CSR min	0.000	max	0.000	mean	0.000		
1493		GPU 64	COO	min	5.070	max	5.140	mean	5.088	1571	bloweya.mtx		H		min	7.205	max	7.205	mean	7.205
1494				CSR min	6.780	max	8.700	mean	7.268	1572	Regular									
1495		H		min	10.572	max	10.584	mean	10.580	1573		GPU 64	COO	min	0.000	max	0.000	mean	0.000	
1496	Column-Gradient					1574				1575			CSR min	0.000	max	0.000	mean	0.000		
1497		GPU 64	COO	min	4.980	max	5.030	mean	5.010	1576		H		min	11.025	max	11.031	mean	11.028	
1498				CSR min	6.390	max	7.640	mean	6.982	1577	Row-Premute		GPU 64	COO	min	4.020	max	4.030	mean	4.023
1499		H		min	10.825	max	10.845	mean	10.836	1578				CSR min	6.070	max	6.750	mean	6.340	
1500	Row-Column-Permute					1579				1580	Row-Gradient		H		min	11.025	max	11.033	mean	11.028
1501		GPU 64	COO	min	4.990	max	5.010	mean	4.997	1581		GPU 64	COO	min	4.090	max	4.160	mean	4.111	
1502				CSR min	6.300	max	7.160	mean	6.636	1582			CSR min	5.980	max	7.370	mean	6.678		
1503		H		min	10.738	max	10.743	mean	10.740	1583		H		min	10.295	max	10.304	mean	10.300	
1504	mult_dcop_01.mtx					1584	Column-Gradient			1585		GPU 64	COO	min	3.980	max	4.010	mean	3.995	
1505	Regular					1586				1587			CSR min	5.880	max	6.780	mean	6.295		
1506		GPU 64	COO	min	5.120	max	5.140	mean	5.134	1588	Row-Column-Permute		H		min	10.881	max	10.887	mean	10.883
1507				CSR min	6.990	max	9.230	mean	7.546	1589		GPU 64	COO	min	4.020	max	4.030	mean	4.023	
1508		H		min	9.689	max	9.689	mean	9.689	1590			CSR min	5.970	max	6.420	mean	6.183		
1509	Row-Premute					1591				1592	lp_osa_07.mtx		H		min	11.025	max	11.033	mean	11.028
1510		GPU 64	COO	min	4.990	max	5.020	mean	5.004	1593	Regular									
1511				CSR min	6.370	max	7.220	mean	6.771	1594		GPU 64	COO	min	4.260	max	4.270	mean	4.261	
1512		H		min	10.738	max	10.743	mean	10.740	1595			CSR min	6.440	max	7.640	mean	6.863		
1513	Row-Gradient					1596				1597	Row-Premute		H		min	8.412	max	8.412	mean	8.412
1514		GPU 64	COO	min	5.060	max	5.100	mean	5.082	1598										
1515				CSR min	6.730	max	7.720	mean	7.317	1599		GPU 64	COO	min	4.200	max	4.200	mean	4.200	
1516		H		min	10.574	max	10.585	mean	10.580	1599			CSR min	6.020	max	7.030	mean	6.418		
1517	Column-Gradient					1600				1601	Row-Gradient		H		min	9.255	max	9.257	mean	9.256
1518		GPU 64	COO	min	4.980	max	5.100	mean	5.012	1602										
1519				CSR min	6.580	max	7.510	mean	7.054	1603		GPU 64	COO	min	4.170	max	4.190	mean	4.180	
1520		H		min	10.828	max	10.842	mean	10.835	1604			CSR min	5.610	max	7.300	mean	5.988		
1521	Row-Column-Permute					1605				1605	Row-Column-Permute		H		min	9.534	max	9.601	mean	9.585
1522		GPU 64	COO	min	4.970	max	5.000	mean	4.986	1606		GPU 64	COO	min	4.190	max	4.190	mean	4.190	
1523				CSR min	6.390	max	7.050	mean	6.677	1607			CSR min	6.070	max	7.000	mean	6.386		
1524		H		min	10.738	max	10.742	mean	10.740	1608		H		min	9.255	max	9.257	mean	9.256	
1525	mult_dcop_02.mtx					1609	Column-Gradient			1609										
1526	Regular					1610				1610	ex19.mtx									
1527		GPU 64	COO	min	5.120	max	5.140	mean	5.133	1611	Regular		GPU 64	COO	min	6.140	max	6.180	mean	6.159
1528				CSR min	6.950	max	7.590	mean	7.336	1612				CSR min	12.780	max	14.400	mean	13.328	
1529		H		min	9.689	max	9.689	mean	9.689	1613		H		min	8.228	max	8.228	mean	8.228	
1530	Row-Premute					1614				1614	Row-Premute									
1531		GPU 64	COO	min	4.970	max	4.990	mean	4.984	1615		GPU 64	COO	min	5.820	max	5.850	mean	5.833	
1532				CSR min	6.440	max	7.110	mean	6.719	1616			CSR min	9.870	max	11.070	mean	10.372		
1533		H		min	10.738	max	10.742	mean	10.740	1617		H		min	11.836	max	11.840	mean	11.838	
1534	Row-Gradient					1618				1618	Row-Gradient									
1535		GPU 64	COO	min	5.070	max	5.150	mean	5.086	1619		GPU 64	COO	min	6.070	max	6.120	mean	6.104	
1536				CSR min	6.650	max	7.930	mean	7.304	1620			CSR min	11.290	max	12.760	mean	12.088		
1537		H		min	10.574	max	10.587	mean	10.580	1621		H		min	10.743	max	10.752	mean	10.748	
1538	Column-Gradient					1622	Column-Gradient													
1539		GPU 64	COO	min	4.980	max	5.040	mean	5.012											
1540				CSR min	6.520	max	7.650	mean	7.139											
1541		H		min	10.829	max	10.846	mean	10.836											
1542	Row-Column-Permute																			
1543		GPU 64	COO	min	4.970	max	5.050	mean	4.983											
1544				CSR min	6.440	max	7.380	mean	6.779											
1545		H		min	10.738	max	10.743	mean	10.740											
1546	lp_fit2d.mtx																			
1547	Regular																			
1548		GPU 64	COO	min	3.960	max	3.960	mean	3.960											

1623		GPU 64 COO min 5.760 max 5.840 mean 5.813	1697		H min 7.380 max 7.380 mean 7.380
1624		CSR min 9.710 max 14.220 mean 10.376	1698	Row-Premute	
1625		H min 11.873 max*11.882 mean 11.878	1699		GPU 64 COO min 4.130 max 4.170 mean 4.134
1626	Row-Column-Permute		1700		CSR min 6.180 max* 9.200 mean 6.796
1627		GPU 64 COO min 5.810 max 5.860 mean 5.838	1701		H min 10.041 max 10.046 mean 10.044
1628		CSR min 9.920 max 10.820 mean 10.240	1702	Row-Gradient	
1629		H min 11.836 max 11.841 mean 11.838	1703		GPU 64 COO min 4.150 max* 4.220 mean 4.163
1630	brainpc2.mtx		1704		CSR min 6.410 max 7.500 mean 6.816
1631	Regular		1705		H min 9.682 max 9.706 mean 9.693
1632		GPU 64 COO min 0.000 max 0.000 mean 0.000	1706	Column-Gradient	
1633		CSR min 0.000 max 0.000 mean 0.000	1707		GPU 64 COO min 4.080 max 4.110 mean 4.096
1634		H min 7.478 max 7.478 mean 7.478	1708		CSR min 6.020 max 7.220 mean 6.309
1635	Row-Premute		1709		H min 10.597 max*10.658 mean 10.631
1636		GPU 64 COO min 4.760 max 4.790 mean 4.773	1710	Row-Column-Permute	
1637		CSR min 6.930 max 7.780 mean 7.310	1711		GPU 64 COO min 4.120 max 4.140 mean 4.130
1638		H min 9.810 max 9.813 mean 9.811	1712		CSR min 6.210 max 7.200 mean 6.609
1639	Row-Gradient		1713		H min 10.041 max 10.046 mean 10.044
1640		GPU 64 COO min 4.820 max* 4.840 mean 4.831	1714	TSOPF_FS_b9_c6.mtx	
1641		CSR min 7.220 max 8.290 mean 7.583	1715	Regular	
1642		H min 9.721 max 9.725 mean 9.723	1716		GPU 64 COO min 0.000 max 0.000 mean 0.000
1643	Column-Gradient		1717		CSR min 0.000 max 0.000 mean 0.000
1644		GPU 64 COO min 4.760 max 4.820 mean 4.779	1718		H min 7.380 max 7.380 mean 7.380
1645		CSR min 6.870 max* 8.300 mean 7.393	1719	Row-Premute	
1646		H min 10.368 max*10.373 mean 10.370	1720		GPU 64 COO min 4.120 max 4.140 mean 4.129
1647	Row-Column-Permute		1721		CSR min 6.170 max 7.160 mean 6.664
1648		GPU 64 COO min 4.750 max 4.780 mean 4.765	1722		H min 10.041 max 10.045 mean 10.043
1649		CSR min 6.940 max 7.580 mean 7.298	1723	Row-Gradient	
1650		H min 9.809 max 9.814 mean 9.811	1724		GPU 64 COO min 4.150 max* 4.180 mean 4.162
1651	shermanACb.mtx		1725		CSR min 6.420 max 7.360 mean 6.723
1652	Regular		1726		H min 9.682 max 9.706 mean 9.693
1653		GPU 64 COO min 4.090 max* 4.130 mean 4.112	1727	Column-Gradient	
1654		CSR min 6.320 max* 7.200 mean 6.779	1728		GPU 64 COO min 4.080 max 4.120 mean 4.096
1655		H min 8.600 max 8.600 mean 8.600	1729		CSR min 5.880 max 7.090 mean 6.403
1656	Row-Premute		1730		H min 10.611 max*10.660 mean 10.637
1657		GPU 64 COO min 4.020 max 4.050 mean 4.036	1731	Row-Column-Permute	
1658		CSR min 5.670 max 6.460 mean 6.014	1732		GPU 64 COO min 4.130 max 4.140 mean 4.130
1659		H min 10.376 max 10.382 mean 10.379	1733		CSR min 6.330 max* 7.390 mean 6.695
1660	Row-Gradient		1734		H min 10.042 max 10.047 mean 10.044
1661		GPU 64 COO min 4.050 max 4.100 mean 4.074	1735	OPF_6000.mtx	
1662		CSR min 5.580 max 6.420 mean 5.996	1736	Regular	
1663		H min 9.918 max 9.924 mean 9.921	1737		GPU 64 COO min 7.270 max* 7.370 mean 7.293
1664	Column-Gradient		1738		CSR min 12.890 max*14.500 mean 13.566
1665		GPU 64 COO min 4.010 max 4.080 mean 4.033	1739		H min 8.799 max 8.799 mean 8.799
1666		CSR min 0.000 max 6.320 mean 5.527	1740	Row-Premute	
1667		H min 10.543 max*10.595 mean 10.589	1741		GPU 64 COO min 6.640 max 6.720 mean 6.678
1668	Row-Column-Permute		1742		CSR min 9.680 max 11.600 mean 10.040
1669		GPU 64 COO min 4.020 max 4.050 mean 4.036	1743		H min 11.873 max 11.877 mean 11.875
1670		CSR min 5.670 max 6.510 mean 6.092	1744	Row-Gradient	
1671		H min 10.377 max 10.381 mean 10.379	1745		GPU 64 COO min 7.090 max 7.140 mean 7.122
1672	cvxqp3.mtx		1746		CSR min 11.250 max 13.030 mean 12.142
1673	Regular		1747		H min 11.110 max 11.117 mean 11.114
1674		GPU 64 COO min 3.500 max* 3.540 mean 3.501	1748	Column-Gradient	
1675		CSR min 11.860 max*13.100 mean 12.694	1749		GPU 64 COO min 6.590 max 6.710 mean 6.644
1676		H min 8.646 max 8.646 mean 8.646	1750		CSR min 9.400 max 13.140 mean 9.991
1677	Row-Premute		1751		H min 12.040 max*12.046 mean 12.043
1678		GPU 64 COO min 3.360 max 3.370 mean 3.365	1752	Row-Column-Permute	
1679		CSR min 6.210 max 7.610 mean 6.631	1753		GPU 64 COO min 6.640 max 6.710 mean 6.679
1680		H min 11.027 max 11.032 mean 11.030	1754		CSR min 9.690 max 10.740 mean 10.050
1681	Row-Gradient		1755		H min 11.874 max 11.877 mean 11.875
1682		GPU 64 COO min 3.370 max 3.380 mean 3.376	1756	OPF_3754.mtx	
1683		CSR min 6.170 max 7.070 mean 6.499	1757	Regular	
1684		H min 11.059 max 11.068 mean 11.064	1758		GPU 64 COO min 4.430 max* 4.450 mean 4.443
1685	Column-Gradient		1759		CSR min 9.710 max*13.000 mean 11.377
1686		GPU 64 COO min 3.350 max 3.390 mean 3.371	1760		H min 8.393 max 8.393 mean 8.393
1687		CSR min 6.150 max 7.180 mean 6.531	1761	Row-Premute	
1688		H min 11.125 max*11.133 mean 11.130	1762		GPU 64 COO min 4.230 max 4.250 mean 4.240
1689	Row-Column-Permute		1763		CSR min 7.430 max 8.750 mean 7.986
1690		GPU 64 COO min 3.350 max 3.380 mean 3.364	1764		H min 11.266 max 11.272 mean 11.269
1691		CSR min 6.040 max 7.440 mean 6.603	1765	Row-Gradient	
1692		H min 11.028 max 11.033 mean 11.030	1766		GPU 64 COO min 4.370 max 4.420 mean 4.382
1693	case9.mtx		1767		CSR min 8.160 max 9.470 mean 8.682
1694	Regular		1768		H min 10.462 max 10.473 mean 10.468
1695		GPU 64 COO min 0.000 max 0.000 mean 0.000	1769	Column-Gradient	
1696		CSR min 0.000 max 0.000 mean 0.000	1770		GPU 64 COO min 4.210 max 4.240 mean 4.227

1771		CSR min 7.160 max 8.080 mean 7.595	1845	Row-Premute	
1772		H min 11.394 max*11.401 mean 11.398	1846		GPU 64 COO min 10.340 max 10.430 mean 10.362
1773	Row-Column-Permute		1847		CSR min 12.880 max 13.340 mean 13.057
1774		GPU 64 COO min 4.230 max 4.250 mean 4.243	1848		H min 10.777 max 10.778 mean 10.777
1775		CSR min 7.230 max 8.940 mean 8.056	1849	Row-Gradient	
1776		H min 11.264 max 11.271 mean 11.269	1850		GPU 64 COO min 10.650 max*10.740 mean 10.688
1777	c-47.mtx		1851		CSR min 12.310 max 13.670 mean 12.562
1778	Regular		1852		H min 11.247 max 11.300 mean 11.281
1779		GPU 64 COO min 5.320 max* 5.340 mean 5.329	1853	Column-Gradient	
1780		CSR min 8.890 max* 9.590 mean 9.249	1854		GPU 64 COO min 10.340 max 10.440 mean 10.398
1781		H min 8.364 max 8.364 mean 8.364	1855		CSR min 9.480 max 10.110 mean 9.782
1782	Row-Premute		1856		H min 12.023 max*12.069 mean 12.047
1783		GPU 64 COO min 5.240 max 5.250 mean 5.241	1857	Row-Column-Permute	
1784		CSR min 7.790 max 8.890 mean 8.214	1858		GPU 64 COO min 10.330 max 10.380 mean 10.356
1785		H min 10.059 max 10.063 mean 10.061	1859		CSR min 12.840 max 13.530 mean 13.119
1786	Row-Gradient		1860		H min 10.776 max 10.778 mean 10.777
1787		GPU 64 COO min 5.230 max 5.260 mean 5.242	1861	aft01.mtx	
1788		CSR min 7.080 max 8.050 mean 7.673	1862	Regular	
1789		H min 10.206 max 10.226 mean 10.218	1863		GPU 64 COO min 3.680 max* 3.690 mean 3.688
1790	Column-Gradient		1864		CSR min 13.860 max*14.830 mean 14.560
1791		GPU 64 COO min 5.080 max 5.120 mean 5.105	1865		H min 7.811 max 7.811 mean 7.811
1792		CSR min 5.780 max 6.970 mean 6.359	1866	Row-Premute	
1793		H min 11.205 max*11.233 mean 11.222	1867		GPU 64 COO min 3.510 max 3.530 mean 3.513
1794	Row-Column-Permute		1868		CSR min 6.420 max 10.520 mean 7.265
1795		GPU 64 COO min 5.220 max 5.250 mean 5.227	1869		H min 11.161 max*11.170 mean 11.165
1796		CSR min 7.860 max 8.710 mean 8.247	1870	Row-Gradient	
1797		H min 10.059 max 10.064 mean 10.061	1871		GPU 64 COO min 3.630 max 3.670 mean 3.643
1798	mhd4800a.mtx		1872		CSR min 10.760 max 13.510 mean 12.199
1799	Regular		1873		H min 10.248 max 10.265 mean 10.258
1800		GPU 64 COO min 3.090 max* 3.100 mean 3.098	1874	Column-Gradient	
1801		CSR min 11.570 max*12.290 mean 12.092	1875		GPU 64 COO min 3.510 max 3.520 mean 3.519
1802		H min 7.132 max 7.132 mean 7.132	1876		CSR min 6.490 max 11.230 mean 7.645
1803	Row-Premute		1877		H min 11.112 max 11.121 mean 11.117
1804		GPU 64 COO min 3.020 max 3.020 mean 3.020	1878	Row-Column-Permute	
1805		CSR min 5.560 max 7.270 mean 6.007	1879		GPU 64 COO min 3.510 max 3.540 mean 3.515
1806		H min 10.959 max*10.968 mean 10.963	1880		CSR min 6.510 max 11.650 mean 7.311
1807	Row-Gradient		1881		H min 11.161 max 11.168 mean 11.165
1808		GPU 64 COO min 3.080 max 3.100 mean 3.088	1882	TSOPF_RS_b39_c7.mtx	
1809		CSR min 10.250 max 12.150 mean 11.340	1883	Regular	
1810		H min 9.509 max 9.528 mean 9.520	1884		GPU 64 COO min 5.970 max* 6.010 mean 5.988
1811	Column-Gradient		1885		CSR min 12.470 max*21.120 mean 13.816
1812		GPU 64 COO min 3.020 max 3.050 mean 3.026	1886		H min 7.304 max 7.304 mean 7.304
1813		CSR min 5.530 max 10.580 mean 6.432	1887	Row-Premute	
1814		H min 10.933 max 10.946 mean 10.939	1888		GPU 64 COO min 5.840 max 5.870 mean 5.856
1815	Row-Column-Permute		1889		CSR min 10.780 max 15.810 mean 11.425
1816		GPU 64 COO min 3.020 max 3.020 mean 3.020	1890		H min 10.537 max 10.540 mean 10.539
1817		CSR min 5.510 max 6.830 mean 6.136	1891	Row-Gradient	
1818		H min 10.959 max 10.967 mean 10.963	1892		GPU 64 COO min 5.950 max 6.000 mean 5.975
1819	gen4.mtx		1893		CSR min 11.520 max 17.250 mean 12.799
1820	Regular		1894		H min 9.638 max 9.646 mean 9.641
1821		GPU 64 COO min 3.300 max* 3.320 mean 3.308	1895	Column-Gradient	
1822		CSR min 5.250 max 6.340 mean 5.705	1896		GPU 64 COO min 5.790 max 5.860 mean 5.827
1823		H min 9.234 max 9.234 mean 9.234	1897		CSR min 10.500 max 14.080 mean 11.237
1824	Row-Premute		1898		H min 11.128 max*11.223 mean 11.209
1825		GPU 64 COO min 3.290 max 3.310 mean 3.299	1899	Row-Column-Permute	
1826		CSR min 5.190 max 7.420 mean 5.683	1900		GPU 64 COO min 5.850 max 5.870 mean 5.855
1827		H min 10.249 max 10.254 mean 10.252	1901		CSR min 10.790 max 15.250 mean 11.718
1828	Row-Gradient		1902		H min 10.537 max 10.541 mean 10.539
1829		GPU 64 COO min 3.300 max 3.310 mean 3.301	1903	mult_dcop_03.mtx	
1830		CSR min 5.370 max 6.310 mean 5.659	1904	Regular	
1831		H min 9.934 max 9.958 mean 9.948	1905		GPU 64 COO min 5.130 max* 5.220 mean 5.142
1832	Column-Gradient		1906		CSR min 7.250 max* 9.320 mean 7.722
1833		GPU 64 COO min 3.240 max 3.260 mean 3.249	1907		H min 9.689 max 9.689 mean 9.689
1834		CSR min 5.090 max* 8.660 mean 5.546	1908	Row-Premute	
1835		H min 10.853 max*10.873 mean 10.864	1909		GPU 64 COO min 4.980 max 5.030 mean 4.999
1836	Row-Column-Permute		1910		CSR min 6.460 max 8.470 mean 6.950
1837		GPU 64 COO min 3.290 max 3.320 mean 3.296	1911		H min 10.738 max 10.742 mean 10.740
1838		CSR min 5.190 max 7.550 mean 5.659	1912	Row-Gradient	
1839		H min 10.249 max 10.255 mean 10.252	1913		GPU 64 COO min 5.070 max 5.140 mean 5.088
1840	Maragal_6.mtx		1914		CSR min 6.780 max 8.700 mean 7.268
1841	Regular		1915		H min 10.572 max 10.584 mean 10.580
1842		GPU 64 COO min 10.580 max 10.620 mean 10.599	1916	Column-Gradient	
1843		CSR min 15.620 max*16.470 mean 15.832	1917		GPU 64 COO min 4.980 max 5.030 mean 5.010
1844		H min 9.930 max 9.930 mean 9.930	1918		CSR min 6.390 max 7.640 mean 6.982

1919		H	min 10.825 max 10.845 mean 10.836	1974
1920	Row-Column-Permute			1975
1921		GPU 64 COO min 4.990 max 5.010 mean 4.997		1976
1922		CSR min 6.300 max 7.160 mean 6.636		1977
1923		H min 10.738 max 10.743 mean 10.740		1978

REFERENCES

- [1] Hartwig Anzt, Terry Cojean, Chen Yen-Chen, Jack J. Dongarra, Goran Flegar, Pratik Nayak, Stanimire Tomov, Yuhsiang M. Tsai, and Weichung Wang. 2020. Load-balancing Sparse Matrix Vector Product Kernels on GPUs. *ACM Trans. Parallel Comput.* 7, 1 (2020), 2:1–2:26. <https://doi.org/10.1145/3380930>
- [2] Paolo D'Alberto, Chris Drome, and Ali Dasdan. 2012. Non-Parametric Methods Applied to the N-Sample Series Comparison. (2012). arXiv:stat.CO/1205.1880
- [3] Jeremy Fowers, Kalin Ovtcharov, Karin Strauss, Eric S Chung, and Greg Stitt. 2014. A high memory bandwidth fpga accelerator for sparse matrix-vector multiplication. In *IEEE Symposium on FPGAs for Custom Computing Machines (FCCM)*. IEEE, 36–43.
- [4] Paul Grigoras, Pavel Burovskiy, and Wayne Luk. 2016. CASK: Open-Source Custom Architectures for Sparse Kernels. In *Proceedings of the International Symposium on Field Programmable Gate Arrays (FPGA)*. ACM, 179–184.
- [5] B. Huang, J. Gao, and X. Li. 2009. An Empirically Optimized Radix Sort for GPU. In *2009 IEEE International Symposium on Parallel and Distributed Processing with Applications*. 234–241.
- [6] Enver Kayaaslan, Cevdet Aykanat, and Bora Uçar. 2018. 1.5D Parallel Sparse Matrix-Vector Multiply. *SIAM J. Scientific Computing* 40, 1 (2018). <https://doi.org/10.1137/16M1105591>
- [7] Srinidhi Kestur, John D Davis, and Eric S Chung. 2012. Towards a universal FPGA matrix-vector multiplication architecture. In *IEEE Symposium on FPGAs for Custom Computing Machines (FCCM)*. IEEE, 9–16.
- [8] Sicheng Li, Yandan Wang, Wujie Wen, Yu Wang, Yiran Chen, and Hai Li. 2016. A data locality-aware design framework for reconfigurable sparse matrix-vector multiplication kernel. In *Proceedings of the International Conference on Computer-Aided Design (ICCAD)*. IEEE, 1–6.
- [9] X. Li, M. J. Garzaran, and D. Padua. 2004. A dynamically tuned sorting library. In *International Symposium on Code Generation and Optimization, 2004. CGO 2004*. 111–122.
- [10] Krishna K Nagar and Jason D Bakos. 2011. A sparse matrix personality for the convey HC-1. In *IEEE Symposium on FPGAs for Custom Computing Machines (FCCM)*. IEEE, 1–8.
- [11] Brian A. Page and Peter M. Kogge. 2018. Scalability of Hybrid Sparse Matrix Dense Vector (SpMV) Multiplication. In *2018 International Conference on High Performance Computing & Simulation, HPCS 2018, Orleans, France, July 16-20, 2018*. IEEE, 406–414. <https://doi.org/10.1109/HPCS.2018.00072>
- [12] Kevin Townsend and Joseph Zambreno. 2013. Reduce, reuse, recycle (r3): A design methodology for sparse matrix vector multiplication on reconfigurable platforms. In *Proceedings of the International Conference on Application-Specific Systems, Architectures and Processors (ASAP)*. IEEE, 185–191.
- [13] Shijie Zhou, Rajgopal Kannan, Viktor K Prasanna, Guna Seetharaman, and Qing Wu. 2019. HitGraph: High-throughput Graph Processing Framework on FPGA. *IEEE Transactions on Parallel and Distributed Systems* 30, 10 (2019), 2249–2264.
- [14] Ling Zhuo and Viktor K Prasanna. 2005. Sparse matrix-vector multiplication on FPGAs. In *Proceedings of the International Symposium on Field Programmable Gate Arrays (FPGA)*. ACM, 63–74.